
Lecture Notes for Machine Learning in Python

Professor Eric Larson
Visualization and Dimensionality Reduction

Class Logistics and Agenda

- Finish Visualization Demo
- Dimensionality Reduction
 - PCA
 - Sampling
 - Kernel Methods

What did we talk about last time?

Visualization

Matplotlib

Seaborn

Plotly

03.Data Visualization.ipynb

Other Tutorials:

<https://t.co/zNzD8Q8w5E>

<http://stanford.edu/~mwaskom/software/seaborn/index.html>


<http://pandas.pydata.org/pandas-docs/stable/visualization.html>

<http://matplotlib.org/examples/index.html>

http://nbviewer.ipynb.org/github/mwaskom/seaborn/blob/master/examples/plotting_distributions.ipynb




Dimensionality Reduction: PCA



Curse of Dimensionality

- When dimensionality increases, data becomes increasingly sparse in the space that it occupies
- Definitions of density and distance between points, which is critical for clustering and outlier detection, become less meaningful



Dimensionality Reduction

- Purpose:
 - Avoid curse of dimensionality
 - Select subsets of independent features
 - Reduce amount of time and memory required by data mining algorithms
 - Allow data to be more easily visualized
 - May help to eliminate irrelevant features or reduce noise
- Techniques
 - Principle Component Analysis
 - Discriminant Analysis
 - Others: supervised and non-linear techniques

Karl Pearson




I invented PCA...
and *social Darwinism*

1857-1936


Dimensionality Reduction: PCA

- Goal is to find a projection that captures the largest amount of variation in data



Dimensionality Reduction: PCA

- Find the **eigenvectors** of the **covariance** matrix
- keep the “k” **largest** eigenvectors



$E1$	$E2$
0.85	0.85
0.52	-0.52

covariance

37.1	-6.7
-6.7	43.9

	$A1$	$A2$
1	66	33.6
2	54	26.6
3	69	23.3
4	73	28.1
5	61	43.1
6	62	25.6

	$A1$	$A2$
1	1.83	3.55
2	-10.1	-3.45
3	4.83	-6.75
4	8.83	-1.95
5	-3.17	13.05
6	-2.17	-4.45

zero mean

Dimensionality Reduction: PCA

attribute₂


now I can represent data
as if it only had one attribute!!
vector weights are also meaningful...

E1

0.85

0.52

component




This projection is called a “Transform”
known as the **Karhunen-Loève Transform**

Explained Variance

- Each principle component explains a certain amount of variation in the data.
- This explained variation is embedded in the eigenvalues for each eigenvector

$$r_q = \frac{\sum_{j=1}^q \lambda_j}{\sum_{j=1}^p \lambda_j}$$



Dimensionality Reduction: PCA

- Genetic profiles distilled to 2 components





image source: 23 and me

Dimensionality Reduction: PCA

- Need more help with the PCA algorithm (and python)?
 - check out Sebastian Raschka's notebooks:
http://nbviewer.ipython.org/github/rasbt/pattern_classification/blob/master/dimensionality_reduction/projection/principal_component_analysis.ipynb

Or check out PCA for dummies:


<https://georgemdallas.wordpress.com/2013/10/30/principal-component-analysis-4-dummies-eigenvectors-eigenvalues-and-dimension-reduction/>



Demo

Dimension Reduction

PCA
biplots



Other Tutorials:


http://scikit-learn.org/stable/auto_examples/decomposition/plot_pca_vs_lda.html#example-decomposition-plot-pca-vs-lda-py

<http://nbviewer.ipython.org/github/ogrisel/notebooks/blob/master/Labeled%20Faces%20in%20the%20Wild%20recognition.ipynb>

For Next Lecture

- Next Lecture:
 - Kernel Methods
 - Dimension Reduction Demo
 - Crash-course Image Feature Extraction

Lecture Notes for Machine Learning in Python



Professor Eric Larson
Dimensionality and Images

Class Logistics and Agenda


- **Logistics:**

- Next lab due in ~1.5 weeks
- Accessing Videos from Canvas

- **Agenda**

- Randomized
- Kernel Methods
- Common Feature Extraction Methods for Images

Last time it was so linear...



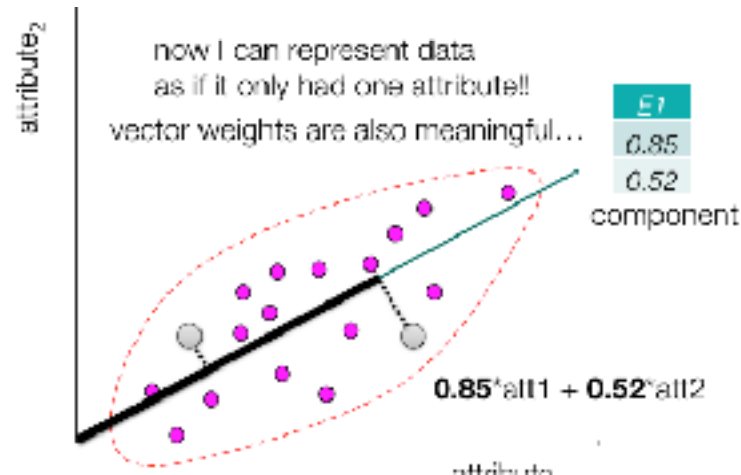
E1	E2
0.85	0.85
0.52	-0.52

37.1	-6.7
-6.7	43.9


	A1	A2
1	66	33.6
2	54	26.6
3	69	23.3
4	73	28.1
5	61	43.1
6	62	25.6

	A1	A2
1	1.83	3.55
2	-10.1	-3.45
3	4.83	-6.75
4	8.83	-1.95
5	-3.17	13.05
6	-2.17	-4.45

zero mean



$$r_q = \frac{\sum_{j=1}^q \lambda_j}{\sum_{j=1}^p \lambda_j}$$



Dimension Reduction

PCA
biplots

Last Time: Demo



Other Tutorials:

http://scikit-learn.org/stable/auto_examples/decomposition/plot_pca_vs_lda.html#example-decomposition-plot-pca-vs-lda-py

<http://nbviewer.ipython.org/github/ogrisel/notebooks/blob/master/Labeled%20Faces%20in%20the%20Wild%20recognition.ipynb>

Self Test ML2b.1

Principal Components Analysis works well for categorical data by design.

- A. True
- B. False
- C. It doesn't but people do it anyway


Dimensionality Reduction: Randomized PCA

- **Problem:** PCA on all that data can take a while to compute
 - What if the number of instances is gigantic?
 - What if the number of dimensions is gigantic?
- What if we partially construct the covariance matrix with a lower rank matrix?
 - By **randomly sampling from the dataset and projecting**, we can get something representative of covariance matrix, but with **lower rank**
 - Gives a matrix with typically good enough precision of actual eigenvectors



$$\|A - QQ^*A\| \leq \left[1 + 11\sqrt{k+p} \cdot \sqrt{\min\{m,n\}}\right] \sigma_{k+1}$$

source: Halko, et al., (2009) Finding structure with randomness: Stochastic algorithms for constructing approximate matrix decompositions. <https://arxiv.org/pdf/0909.4061.pdf>

Non-linear Dimensionality Reduction




Dimensionality Reduction: non-linear



- Sometimes a **linear transform** is not enough
- A powerful non-linear transform has seen a resurgence in past decade: **kernel PCA**

Kernel PCA

- Estimate Covariance in higher dimensional space
- Get eigen vectors from nonlinear dot product
- Projecting onto these can be understood as principle components from a higher dimensional space



$\phi(A1) \cdot \phi(A1)$	$\phi(A2) \cdot \phi(A1)$
$\phi(A2) \cdot \phi(A1)$	$\phi(A2) \cdot \phi(A2)$

	A1	A2
1	66	33.6
2	54	26.6
3	69	23.3
4	73	28.1
5	61	43.1
6	62	25.6


Kernel PCA

kernel: defines what the dot product is in higher dimensional space

$$\kappa(A1, A2) = \phi(A1) \cdot \phi(A2)$$

some kernels have corresponding transformations with **infinite dimensions!!**

$\phi(A1) \cdot \phi(A1)$	$\phi(A2) \cdot \phi(A1)$
$\phi(A2) \cdot \phi(A1)$	$\phi(A2) \cdot \phi(A2)$



- **Key insight:** don't need to know the actual principle components, just the projections
- **Never need** eigen vectors of **full** covariance matrix, just how much the vectors co-vary in higher space!

	A1	A2
1	66	33.6
2	54	26.6
3	69	23.3
4	73	28.1
5	61	43.1
6	62	25.6

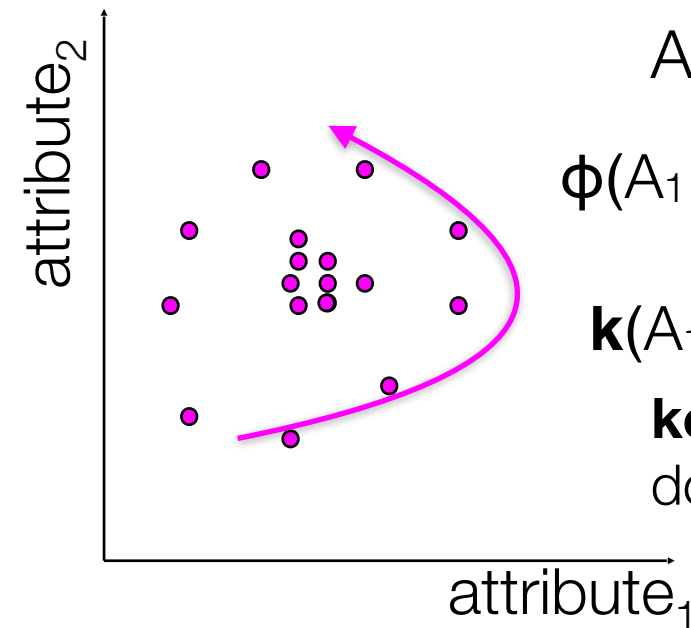
Kernel PCA

kernel: defines what the dot product is in higher dimensional space

some kernels have corresponding transformations with **infinite dimensions!!**

$\phi(A1) \cdot \phi(A1)$	$\phi(A2) \cdot \phi(A1)$
$\phi(A2) \cdot \phi(A1)$	$\phi(A2) \cdot \phi(A2)$

	A1	A2
1	66	33.6
2	54	26.6
3	69	23.3
4	73	28.1
5	61	43.1
6	62	25.6



$$A_1 = [a_1 \ a_2]^T$$

$$\phi(A_1) = [a_1 \ a_2 \ a_1 \cdot a_2 \ a_1^2 \ a_1 \cdot a_2^3 \dots]^T$$

$$\mathbf{k}(A_1, A_2) = \exp(-\gamma \|A_1 - A_2\|^2)$$

kernel: radial basis function (rbf)
dot product in higher dimensional space

Kernel PCA

images: sklearn documentation

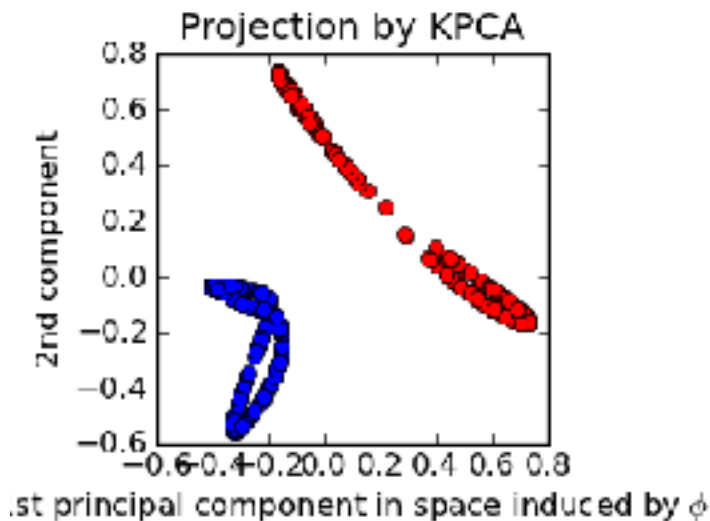

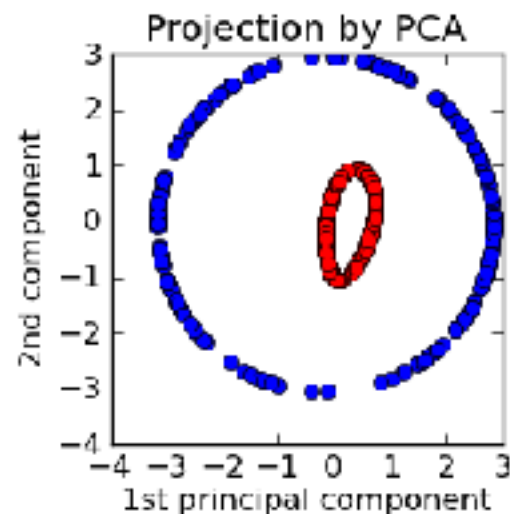
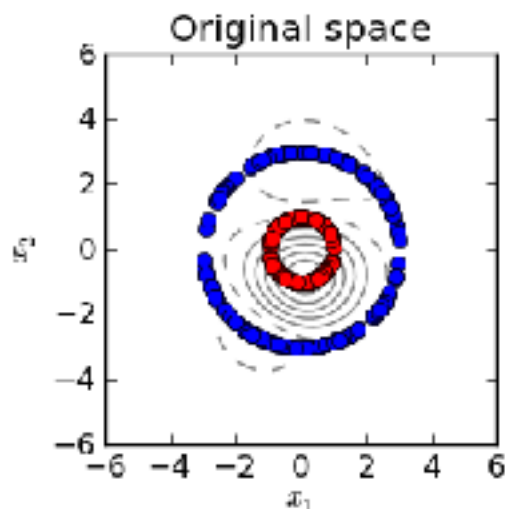




Image Processing and Representation



Images as data

- an image can be represented in many ways
- most common format is a matrix of pixels

- each “pixel” is BGR(A)



- used for capture and display






Image Representation

- need a compact representation

- **grayscale**

$$0.3 \cdot R + 0.59 \cdot G + 0.11 \cdot B,$$

“luminance”

gray

1	4	2	5	6	9
1	4	2	5	5	9
1	4	2	8	8	7
3	4	3	9	9	8
1	0	2	7	7	9
1	4	3	9	8	6
2	4	2	8	7	9

Numpy Matrix
`image[rows, cols]`

on

R

G

B

	1	4	2	5	6	9	
	1	4	2	5	6	9	9
1	4	2	5	6	9	9	7
1	4	2	5	5	9	7	8
1	4	2	8	8	7	8	9
3	4	3	9	9	8	9	6
1	0	2	7	7	9	6	9
1	4	3	9	8	6	9	
2	4	2	8	7	9		

Numpy Matrix
`image[rows, cols, channels]`

Image Representation, Features

Problem: need to represent image as table data

1	4	2	5	6	9
1	4	2	5	5	9
1	4	2	8	8	7
3	4	3	9	9	8
1	0	2	7	7	9
1	4	3	9	8	6
2	4	2	8	7	9

Image Representation, Features

Problem: need to represent image as table data

Solution: row concatenation (also, vectorizing)

Row 1	1	4	2	5	6	9	1	4	2	5	5	9	1	4	2	8	8	7	3	...
Row 2	1	4	2	8	8	7	3	4	3	9	9	8	1	4	2	5	5	9	1	...
...																				
Row N	9	4	6	8	8	7	4	1	3	9	2	1	1	5	2	1	5	9	1	...

Self test: 3a-1

- When vectorizing images into table data, each “feature column” corresponds to:
 - a. the value (color) of pixel
 - b. the spatial location of a pixel in the image
 - c. the size of the image
 - d. the spatial location and color channel of a pixel in an image

Dimension Reduction with Images

Images Representation

Randomized PCA

Kernel PCA



04.Dimension Reduction and Images.ipynb

Features of Images





Image Representation

- need a compact representation

- **hsv**

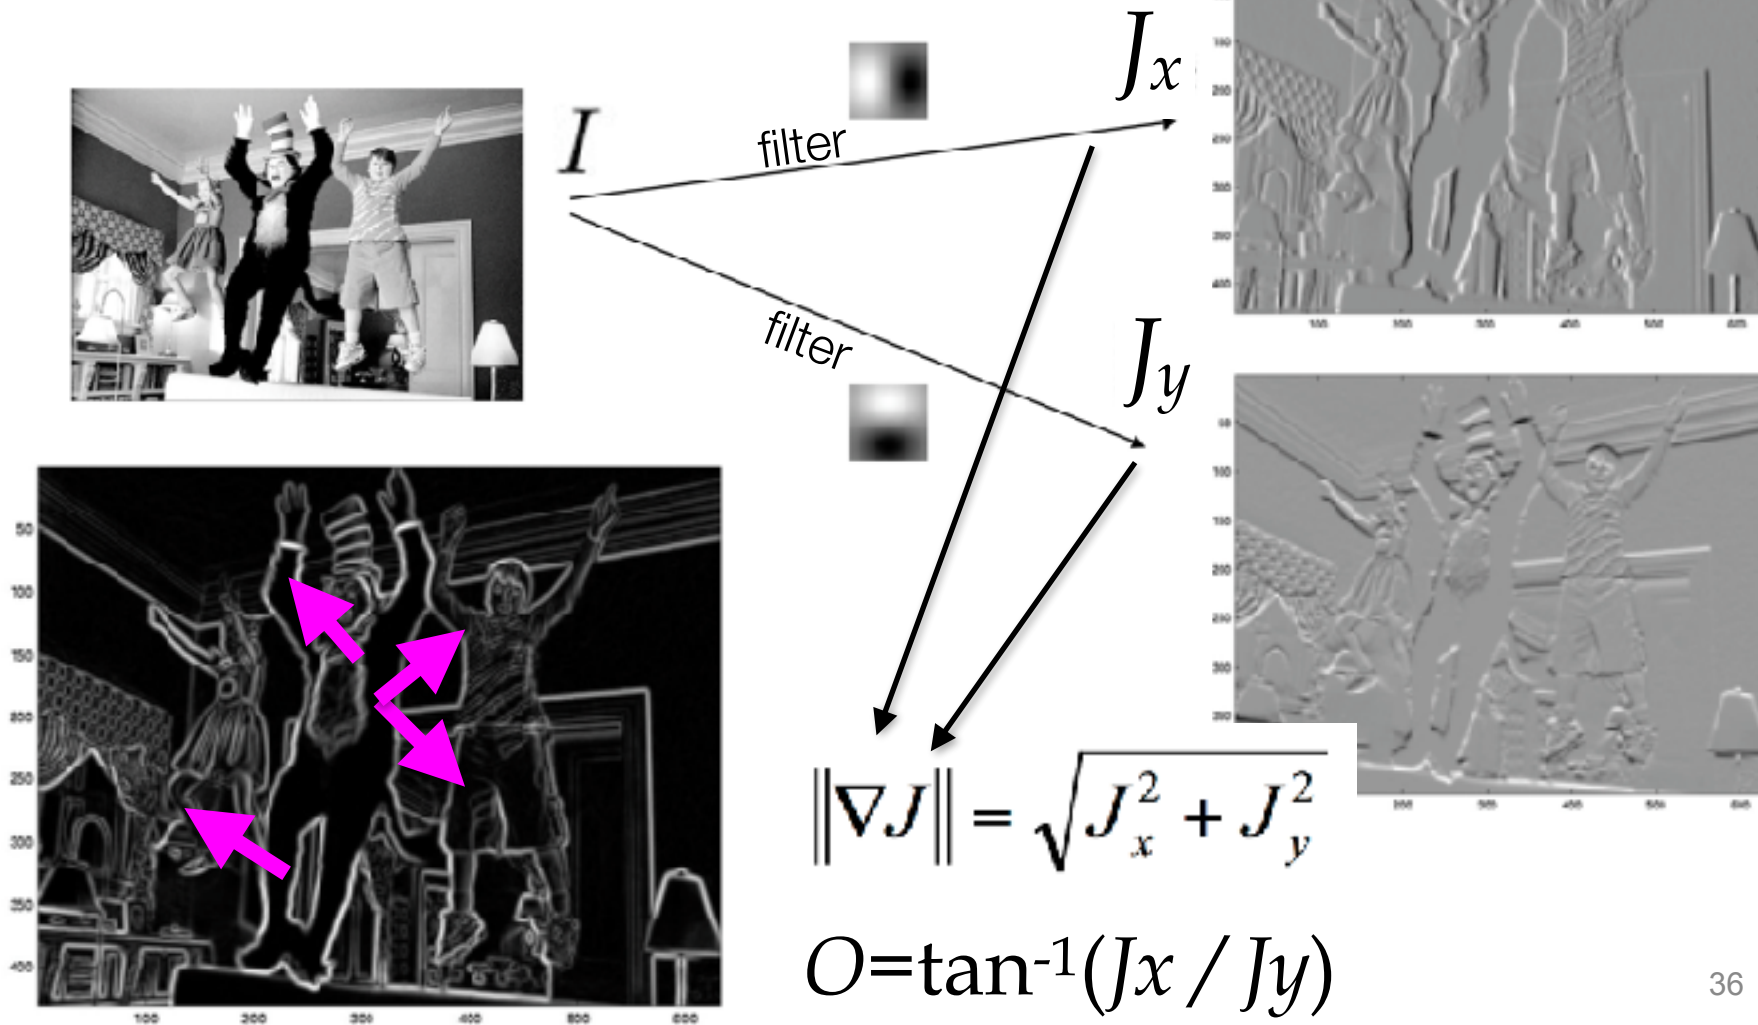
- what we perceive as color (ish)
 - hue: the color value
 - saturation: the richness of the color relative to brightness
 - value: the gray level intensity




R							
G							
B							
1	1	2	5	6	0		
1	1	2	5	6	0	0	0
1	1	2	5	6	0	0	7
1	1	2	5	5	0	7	0
1	1	2	0	0	7	0	0
2	1	2	0	0	0	0	6
V							
1	1	2	5	6	0	6	0
S							
H							
1	1	2	5	6	0	0	0
1	1	2	5	6	0	0	7
1	1	2	5	5	0	7	0
1	1	2	0	0	7	0	0
2	1	2	0	0	0	0	6
1	0	2	7	7	0	6	0
1	1	2	0	0	6	0	
2	1	2	0	7	0		

Common operations

- the gradient (2D derivative)

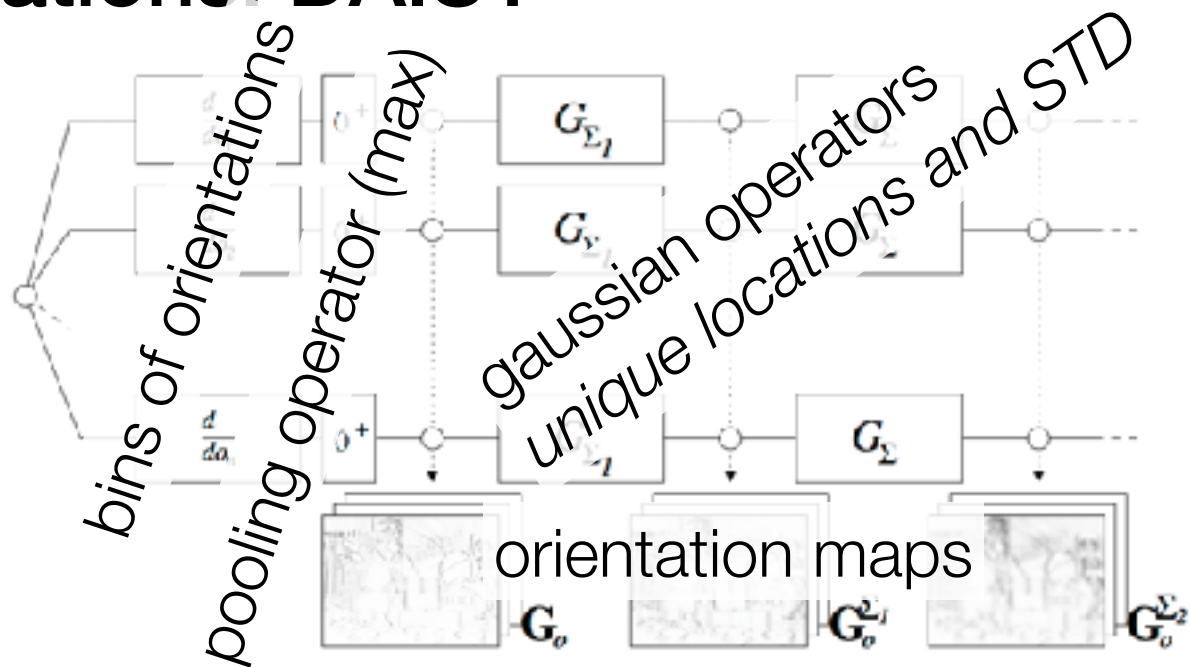


Common operations: DAISY



$$\mathcal{D}(u_0, v_0) =$$

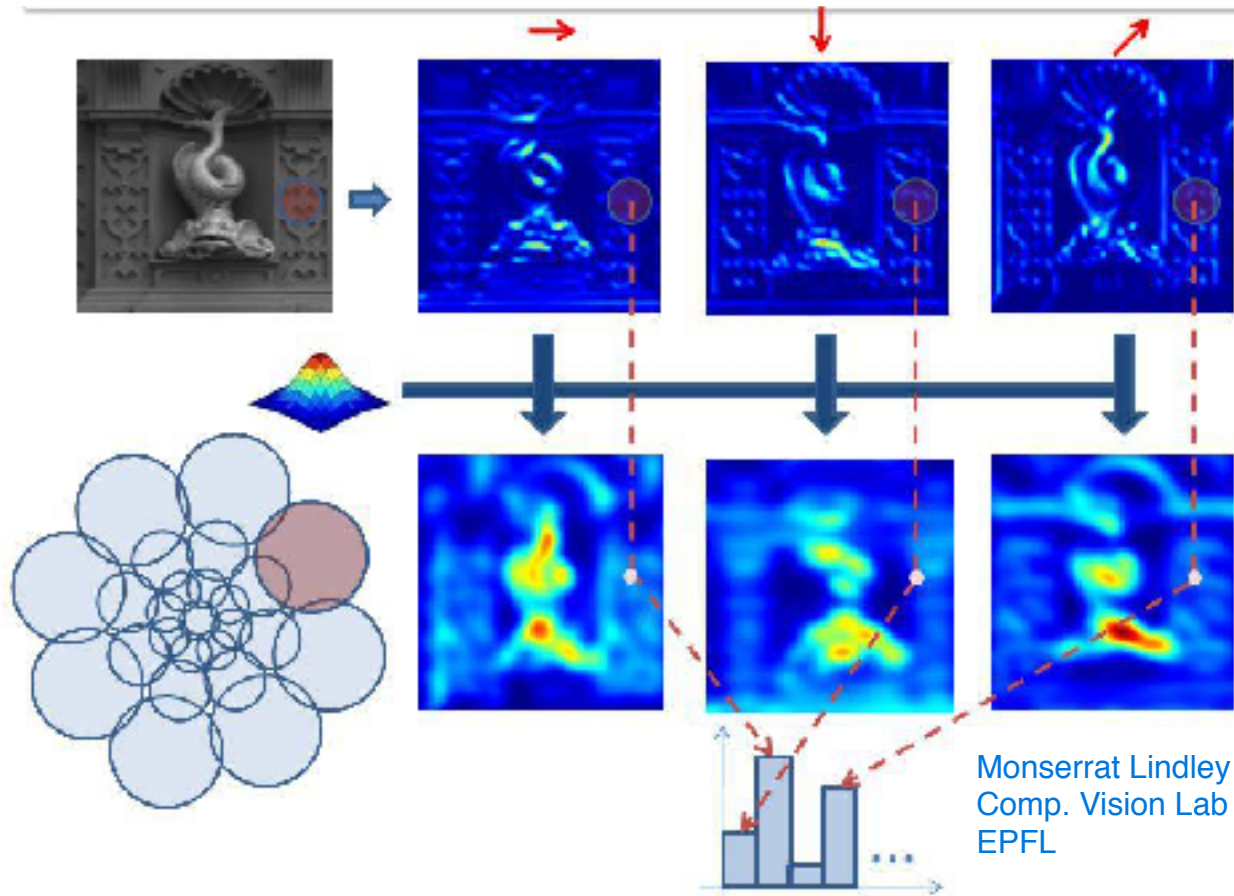

$$\left[\begin{array}{l} \tilde{h}_{\Sigma_1}^\top(u_0, v_0), \\ \tilde{h}_{\Sigma_1}^\top(\mathbf{l}_1(u_0, v_0, R_1)), \dots, \tilde{h}_{\Sigma_1}^\top(\mathbf{l}_T(u_0, v_0, R_1)), \\ \tilde{h}_{\Sigma_2}^\top(\mathbf{l}_1(u_0, v_0, R_2)), \dots, \tilde{h}_{\Sigma_2}^\top(\mathbf{l}_T(u_0, v_0, R_2)), \end{array} \right]$$



take normalized histogram at point u, v

$$\tilde{h}_\Sigma(u, v) = \left\| \left[\mathbf{G}_1^\Sigma(u, v), \dots, \mathbf{G}_H^\Sigma(u, v) \right]^\top \right\|$$

Tola et al. "Daisy: An efficient dense descriptor applied to wide-baseline stereo." Pattern Analysis and Machine Intelligence, IEEE Transactions



take normalized histogram at point u, v


$$\mathcal{D}(u_0, v_0) = \begin{bmatrix} \tilde{\mathbf{h}}_{\Sigma_1}^{\top}(u_0, v_0), \\ \tilde{\mathbf{h}}_{\Sigma_1}^{\top}(\mathbf{l}_1(u_0, v_0, R_1)), \dots, \tilde{\mathbf{h}}_{\Sigma_1}^{\top}(\mathbf{l}_T(u_0, v_0, R_1)), \\ \tilde{\mathbf{h}}_{\Sigma_2}^{\top}(\mathbf{l}_1(u_0, v_0, R_2)), \dots, \tilde{\mathbf{h}}_{\Sigma_2}^{\top}(\mathbf{l}_T(u_0, v_0, R_2)), \end{bmatrix}$$

$$\tilde{\mathbf{h}}_{\Sigma}(u, v) = \left\| \left[\mathbf{G}_1^{\Sigma}(u, v), \dots, \mathbf{G}_H^{\Sigma}(u, v) \right]^{\top} \right\|$$

Tola et al. "Daisy: An efficient dense descriptor applied to wide-baseline stereo." Pattern Analysis and Machine Intelligence, IEEE Transactions

Common operations: DAISY

Num histograms/per ring



Num rings

num orientations making up each histogram

ep




Params:

step, radius, num rings,
num histograms per ring,
orientations per histogram

Common operations: DAISY

Num histograms/per ring



Num rings



Bag of Features

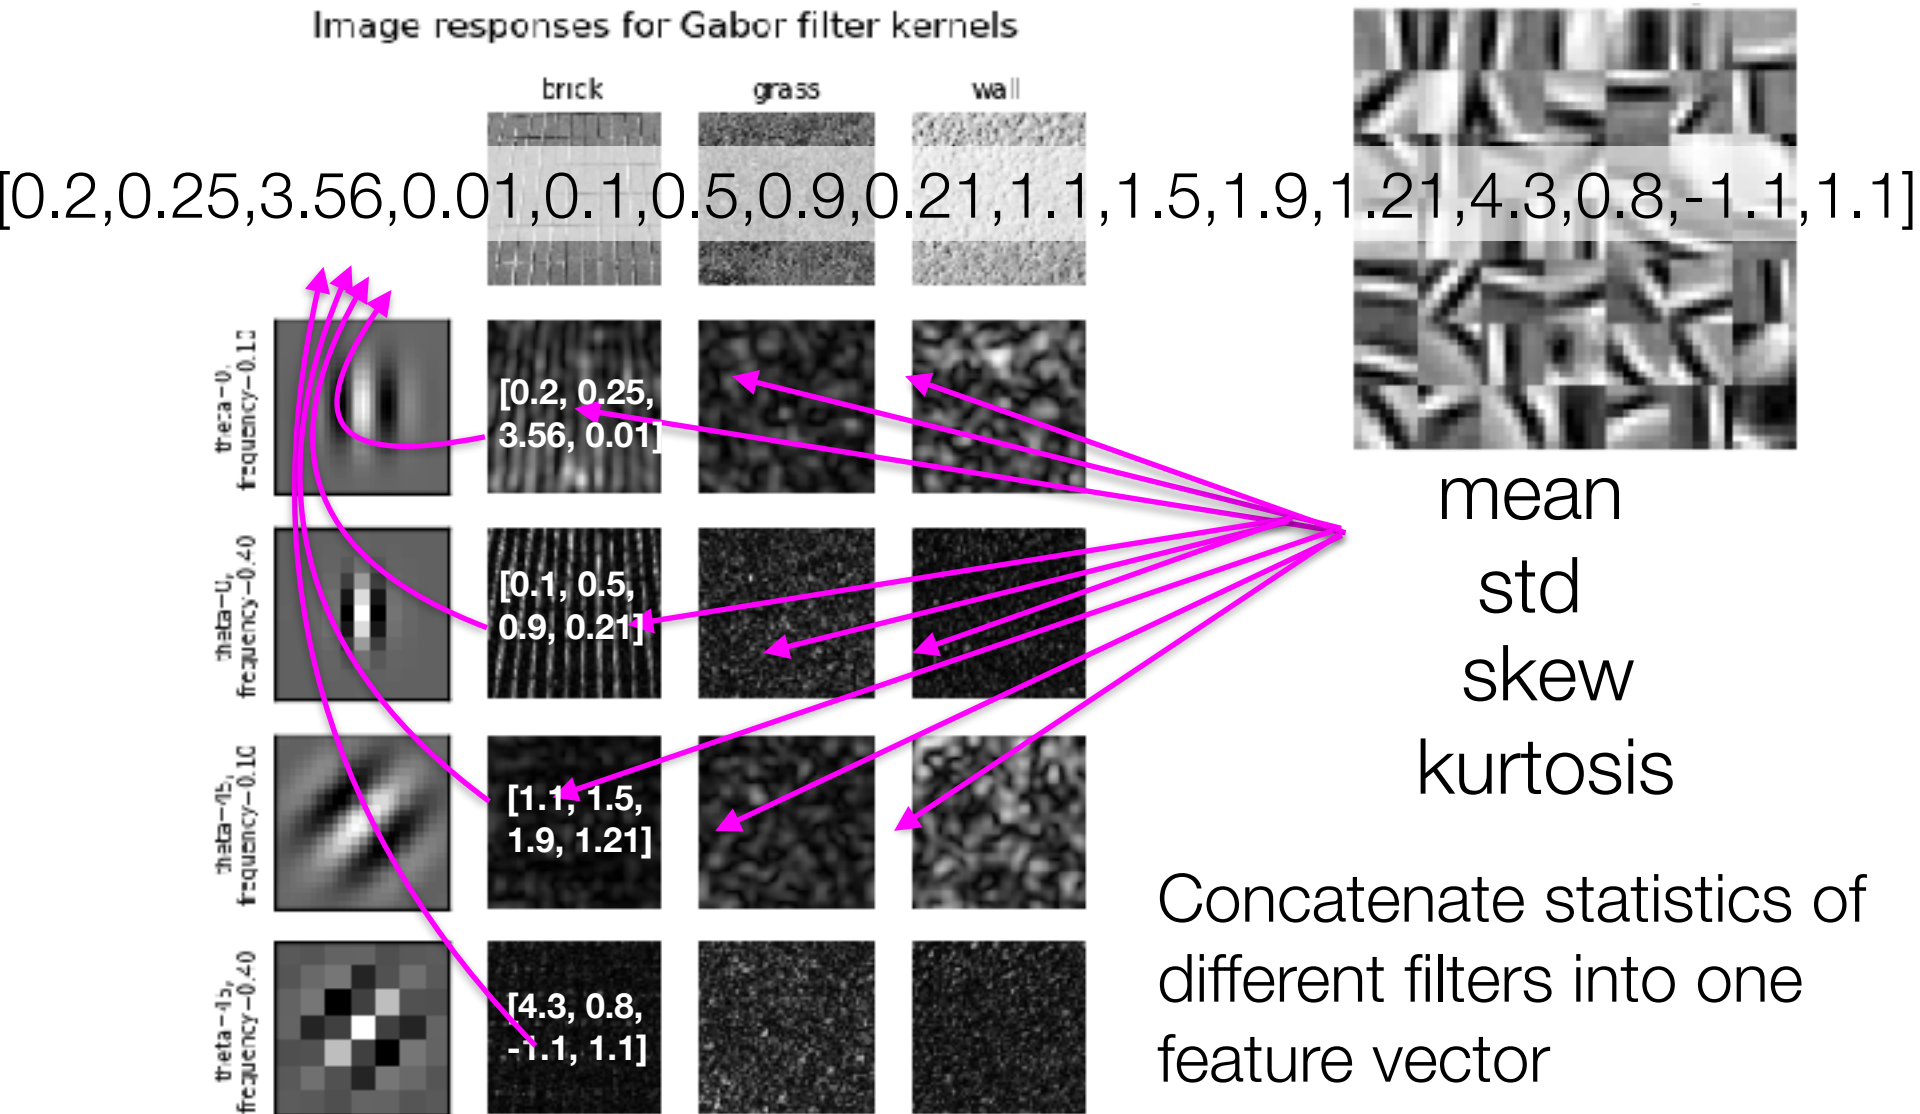
Image Representation

Params:

step, radius, num rings,
num histograms per ring,
orientations per histogram

Common operations: Gabor filter Banks (if time)

- common used for texture classification



More Image Processing

Gradients

DAISY

Gabor Filter Banks



Other Tutorials:

http://scikit-image.org/docs/dev/auto_examples/

For Next Lecture

- Work on your text datasets!
- **Quiz is now live:** Image Processing
- **Next Time:** In-Class Assignment One!!!
- **Next Week:** Project Questions Lecture

Supplemental Slides

- Peruse these at your own leisure!
- These slides might assist you as additional visual aides
- **Slides courtesy of Tan, Steinbach, Kumar**
 - **Introduction to Data Mining**


Dimensionality Reduction: LDA

- PCA tell us variance explained by the data in different directions, but it ignores class labels
- Is there a way to find “components” that will help with **discriminate** between the classes?

$$\arg \max_{comp.} \frac{\sum \text{differences between classes}}{\sum \text{variance within classes}}$$


- called Fisher’s discriminant
- ...but we need to solve this using using *Lagrange multipliers* and gradient-based optimization
- which we haven’t covered yet

I invented Lagrange multipliers... and ...*nothing* impresses me...



Dimensionality Reduction: LDA versus QDA


$$\arg \max_{comp.} \frac{\sum \text{differences between classes}}{\sum \text{variance within classes}}$$



Dimensionality Reduction: LDA versus QDA

$$\arg \max_{comp.} \frac{\sum \text{differences between classes}}{\sum \text{variance within classes}}$$

- “*differences between classes*” is calculated by trying to separate the **mean value** of each **feature** in each **class**
- Linear discriminant analysis:
 - assume the covariance in each class is the same
- Quadrature discriminant analysis:
 - estimate the covariance for each class



Self Test ML2b.2

LDA only allows as many components as there are unique classes in a dataset.

- A. True
- B. False