

Lecture Notes for Machine Learning in Python



Professor Eric Larson
Dimensionality and Images

Class Logistics and Agenda

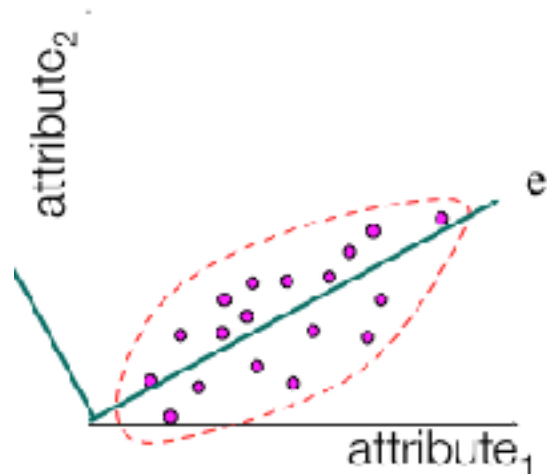
- **Logistics:**

- Next lab due at the end of the week
- Text with visuals!
- Seminar Wednesday on Maneframe
- Next time: no class if we finish everything today
 - Office Hours Location?

- **Agenda**

- Kernel Methods
- Common Feature Extraction Methods for Images

Last time it was so linear...




$E1$	$E2$
0.85	0.85
0.52	-0.52

37.1	-6.7
-6.7	43.9

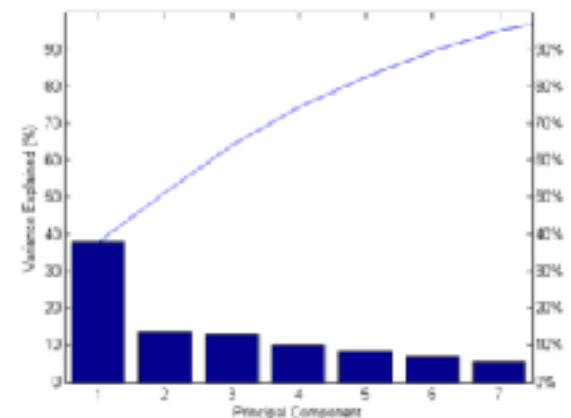
	A1	A2
1	66	33.6
2	54	26.6
3	69	23.3
4	73	28.1
5	61	43.1
6	62	25.6

	A1	A2
1	1.83	3.55
2	-10.1	-3.45
3	4.83	-6.75
4	8.83	-1.95
5	-3.17	13.05
6	-2.17	-4.45

zero mean



$$r_q = \frac{\sum_{j=1}^q \lambda_j}{\sum_{j=1}^p \lambda_j}$$




Dimensionality Reduction: Randomized PCA

- **Problem:** PCA on all that data can take a while to compute
 - What if the number of dimensions is gigantic?
 - Actually, **that's okay**: there are **iterative** algorithms for finding the **largest eigenvalues** that scales well with the number of data dimensions, but **not** the **number of instances...**
 - What if the number of instances is gigantic?
- What if we construct the covariance matrix with a subsample of the data?
 - By **randomly sampling from the dataset**, we can get something representative of the covariance for the entire dataset (if we sample correctly)

Demo

Dimension Reduction

PCA
biplots




Other Tutorials:



http://scikit-learn.org/stable/auto_examples/decomposition/plot_pca_vs_lda.html#example-decomposition-plot-pca-vs-lda-py

<http://nbviewer.ipython.org/github/ogrisel/notebooks/blob/master/Labeled%20Faces%20in%20the%20Wild%20recognition.ipynb>

Non-linear Dimensionality Reduction




Dimensionality Reduction: non-linear



- Sometimes a **linear transform** is not enough
- A powerful non-linear transform has seen a resurgence in past decade: **kernel PCA**

Kernel PCA

- Estimate Covariance in higher dimensional space
- Get eigen vectors from nonlinear dot product
- Projecting onto these can be understood as principle components from a higher dimensional space



$\phi(A1) \cdot \phi(A1)$	$\phi(A2) \cdot \phi(A1)$
$\phi(A2) \cdot \phi(A1)$	$\phi(A2) \cdot \phi(A2)$


	A1	A2
1	66	33.6
2	54	26.6
3	69	23.3
4	73	28.1
5	61	43.1
6	62	25.6

Kernel PCA

kernel: defines what the dot product is in higher dimensional space

some kernels have corresponding transformations with **infinite dimensions!!**

$\phi(A1) \cdot \phi(A1)$	$\phi(A2) \cdot \phi(A1)$
$\phi(A2) \cdot \phi(A1)$	$\phi(A2) \cdot \phi(A2)$



- **Key insight:** don't need to know the actual principle components, just the projections
- **Never need** eigen vectors of full covariance matrix

	A1	A2
1	66	33.6
2	54	26.6
3	69	23.3
4	73	28.1
5	61	43.1
6	62	25.6

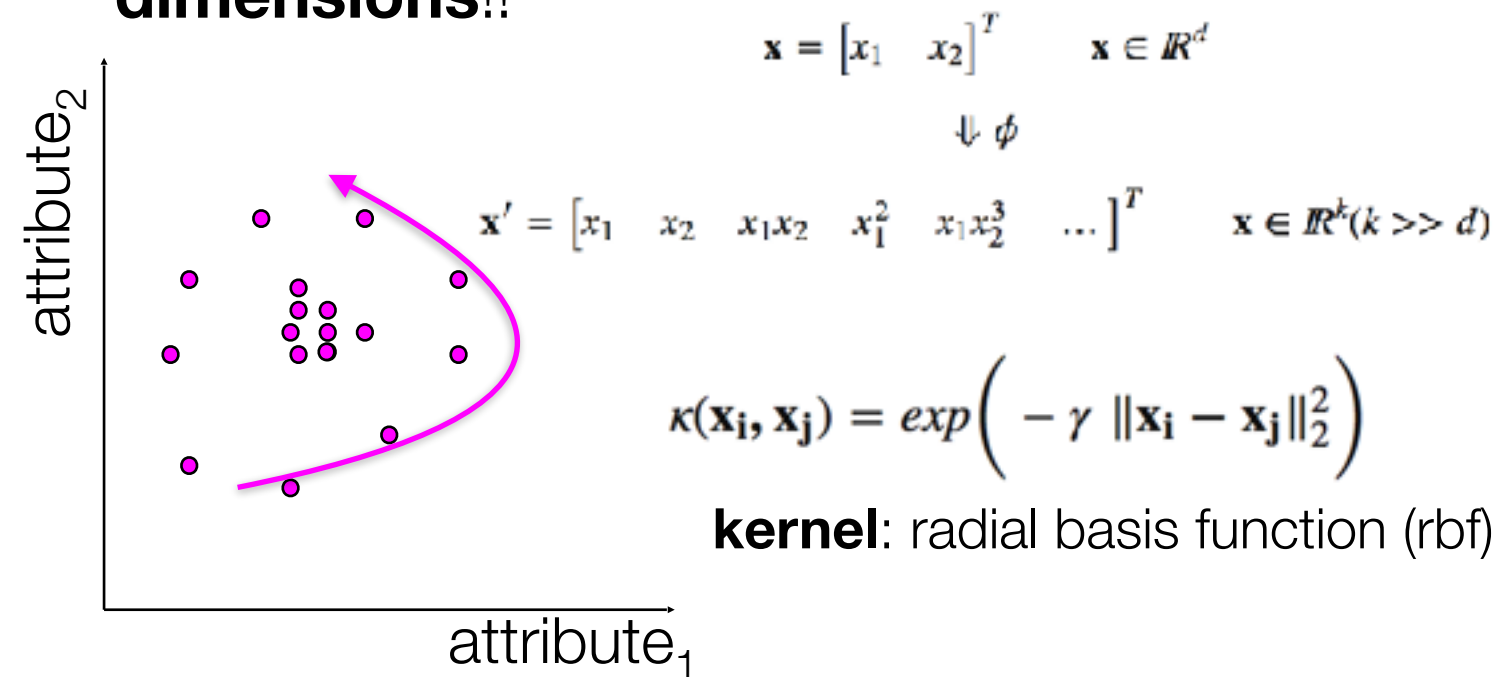
Kernel PCA

kernel: defines what the dot product is in higher dimensional space

some kernels have corresponding transformations with **infinite dimensions!!**

$\phi(A1) \cdot \phi(A1)$	$\phi(A2) \cdot \phi(A1)$
$\phi(A2) \cdot \phi(A1)$	$\phi(A2) \cdot \phi(A2)$

	A1	A2
1	66	33.6
2	54	26.6
3	69	23.3
4	73	28.1
5	61	43.1
6	62	25.6



Kernel PCA





Image Processing and Representation




What is image processing

- the **art** and **science** of manipulating pixels
 - combining images (blending or compositing)
 - enhancing edges and lines
 - adjusting contrast, color
 - warping, transformation
 - filtering
 - features extraction

Images as data

- an image can be represented in many ways
- most common format is a matrix of pixels

- each “pixel” is BGR(A)



- used for capture and display






Image Representation

- need a compact representation

- **grayscale**

$$0.3 \cdot R + 0.59 \cdot G + 0.11 \cdot B,$$

“luminance”

gray

1	4	2	5	6	9
1	4	2	5	5	9
1	4	2	8	8	7
3	4	3	9	9	8
1	0	2	7	7	9
1	4	3	9	8	6
2	4	2	8	7	9

Numpy Matrix
`image[rows, cols]`

						R	
						G	
						B	
						1	4
						2	5
						6	9
						9	9
						9	7
						7	8
						8	9
						9	6
						6	9
						9	

Numpy Matrix
`image[rows, cols, channels]`

Image Representation, Features

Problem: need to represent image as table data

1	4	2	5	6	9
1	4	2	5	5	9
1	4	2	8	8	7
3	4	3	9	9	8
1	0	2	7	7	9
1	4	3	9	8	6
2	4	2	8	7	9

Image Representation, Features

Problem: need to represent image as table data

Solution: row concatenation

Row 1	1	4	2	5	6	9	1	4	2	5	5	9	1	4	2	8	8	7	3	...
Row 2	1	4	2	8	8	7	3	4	3	9	9	8	1	4	2	5	5	9	1	...
...																				
Row N	9	4	6	8	8	7	4	1	3	9	2	1	1	5	2	1	5	9	1	...

Self test: 3a-1

- When vectorizing images into table data, each feature column corresponds to:
 - a. the value (color) of pixel
 - b. the spatial location of a pixel in the image
 - c. the size of the image
 - d. the spatial location and color channel of a pixel in an image

Dimension Reduction with Images

Images Representation

Randomized PCA

Kernel PCA



04.Dimension Reduction and Images.ipynb

Features of Images





Image Representation

- need a compact representation

•hsv

- what we perceive as color (ish)
 - hue: the color value
 - saturation: the richness of the color relative to brightness
 - value: the gray level intensity



on

R

G

B

V

S

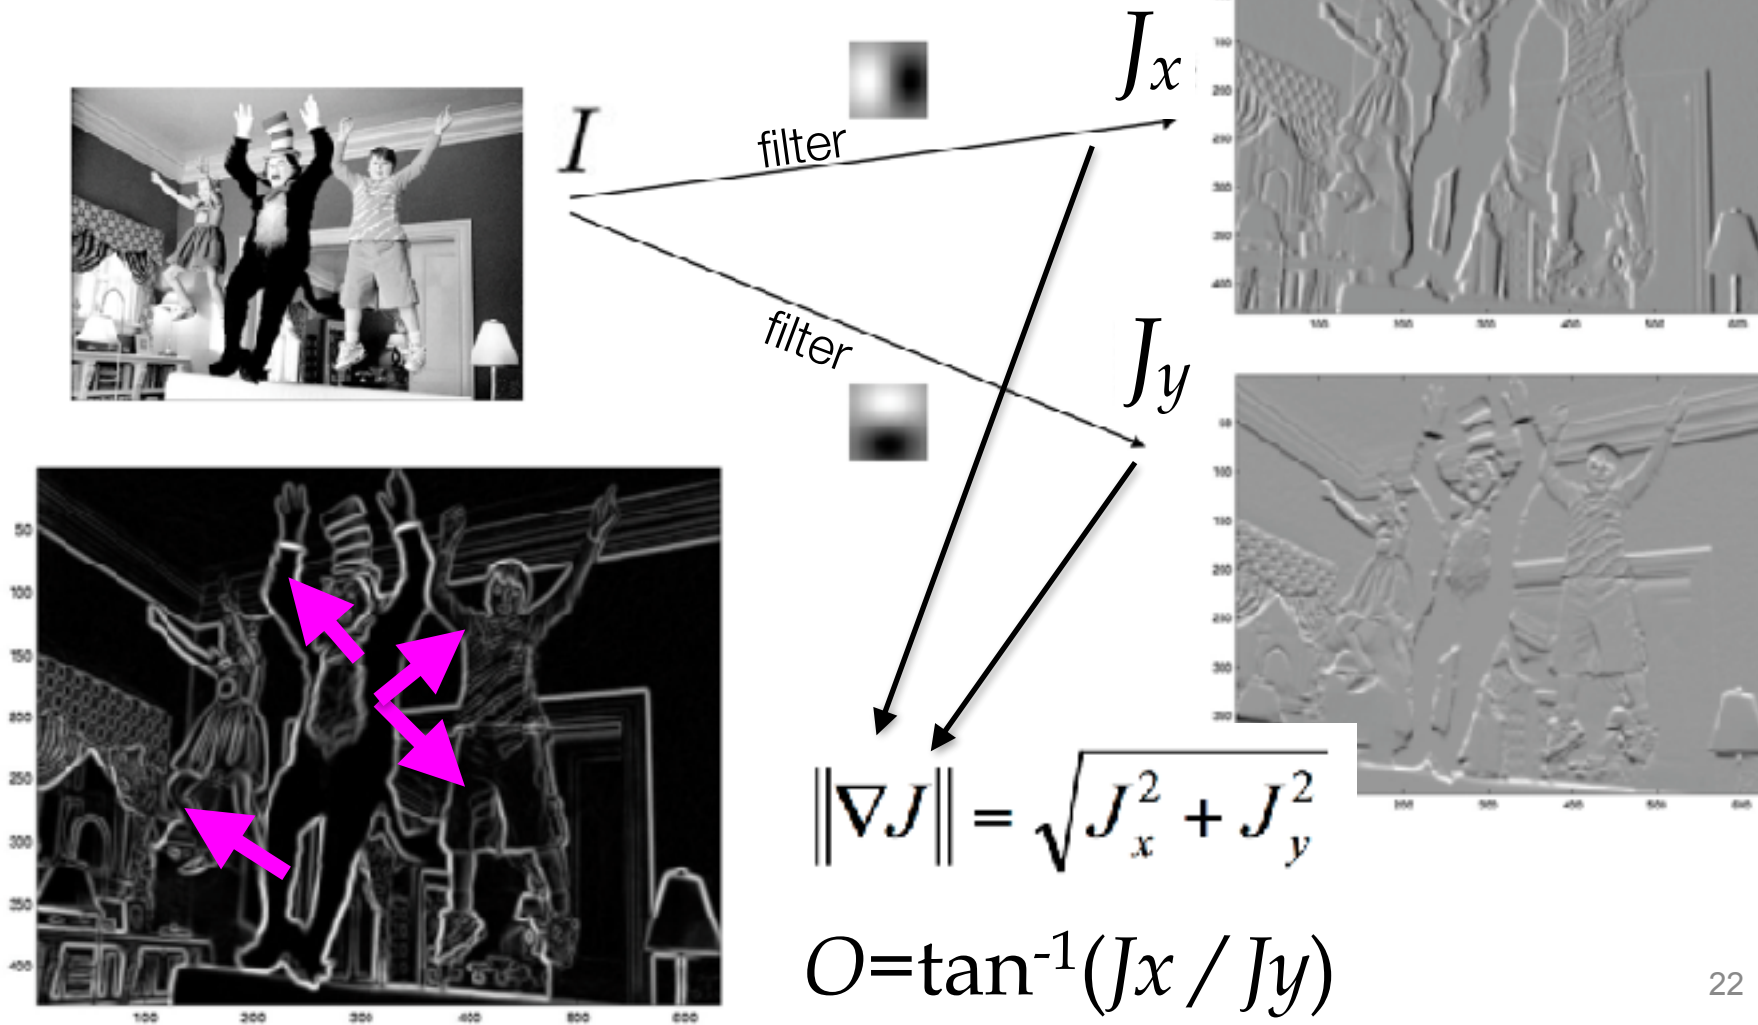
H

			1	1	2	5	6	0	
		1	1	2	5	6	0	0	
	1	1	2	5	6	0	0	7	
	1	1	2	5	5	0	7	0	
	1	1	2	0	0	7	0	0	
	2	1	2	0	0	0	0	6	
	1	0	2	7	7	0	6	0	
	1	1	2	0	0	6	0		
	2	1	2	0	7	0			


21

Common operations

- the gradient (2D derivative)




Common operations: DAISY



$$\mathcal{D}(u_0, v_0) =$$

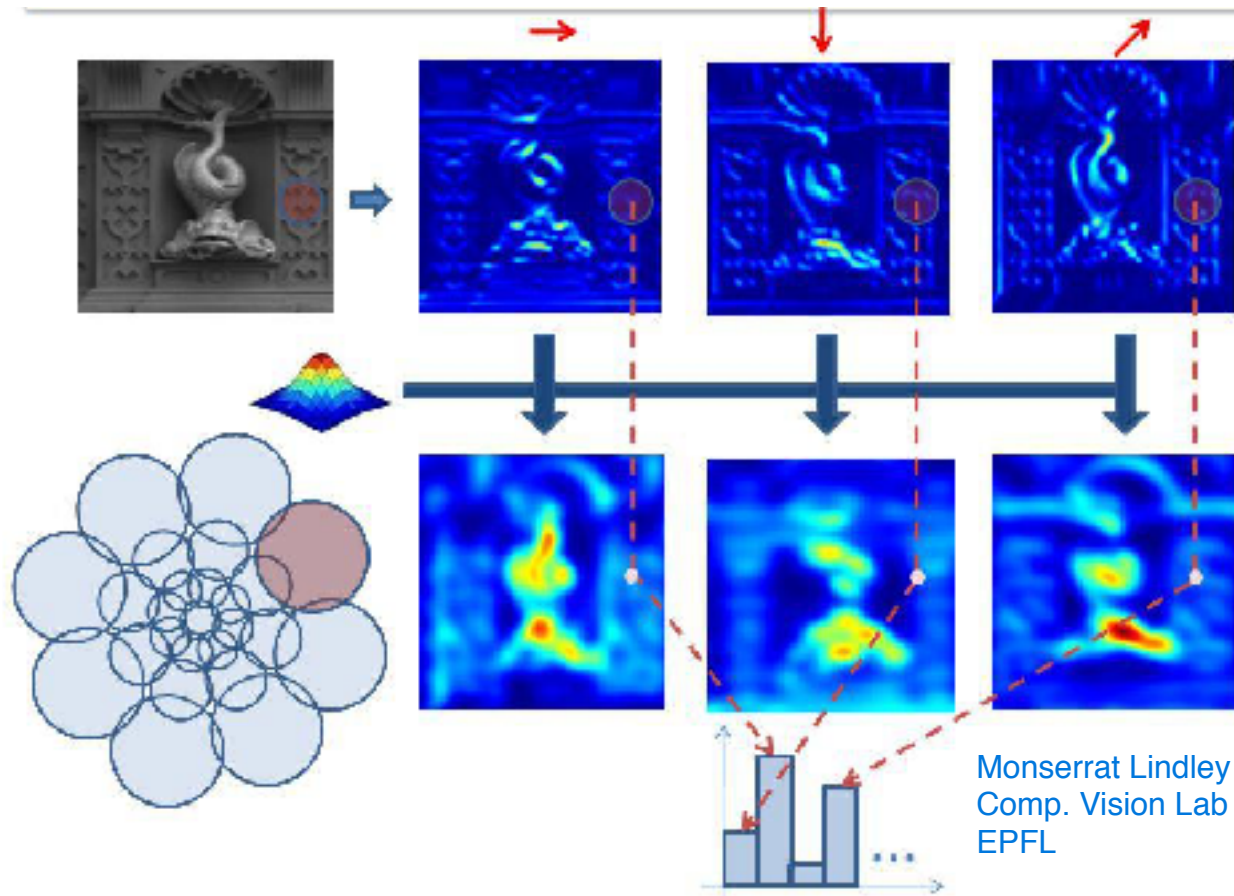

$$\left[\begin{array}{l} \tilde{h}_{\Sigma_1}^\top(u_0, v_0), \\ \tilde{h}_{\Sigma_1}^\top(l_1(u_0, v_0, R_1)), \dots, \tilde{h}_{\Sigma_1}^\top(l_T(u_0, v_0, R_1)), \\ \tilde{h}_{\Sigma_2}^\top(l_1(u_0, v_0, R_2)), \dots, \tilde{h}_{\Sigma_2}^\top(l_T(u_0, v_0, R_2)), \end{array} \right]$$



take normalized histogram at point u, v

$$\tilde{h}_\Sigma(u, v) = \left\| \left[\mathbf{G}_1^\Sigma(u, v), \dots, \mathbf{G}_H^\Sigma(u, v) \right]^\top \right\|$$

Tola et al. "Daisy: An efficient dense descriptor applied to wide-baseline stereo." Pattern Analysis and Machine Intelligence, IEEE Transactions



take normalized histogram at point u, v


$$\mathcal{D}(u_0, v_0) = \begin{bmatrix} \tilde{\mathbf{h}}_{\Sigma_1}^T(u_0, v_0), \\ \tilde{\mathbf{h}}_{\Sigma_1}^T(\mathbf{l}_1(u_0, v_0, R_1)), \dots, \tilde{\mathbf{h}}_{\Sigma_1}^T(\mathbf{l}_T(u_0, v_0, R_1)), \\ \tilde{\mathbf{h}}_{\Sigma_2}^T(\mathbf{l}_1(u_0, v_0, R_2)), \dots, \tilde{\mathbf{h}}_{\Sigma_2}^T(\mathbf{l}_T(u_0, v_0, R_2)), \end{bmatrix}$$

$$\tilde{\mathbf{h}}_{\Sigma}(u, v) = \left\| \left[\mathbf{G}_1^{\Sigma}(u, v), \dots, \mathbf{G}_H^{\Sigma}(u, v) \right]^T \right\|$$

Tola et al. "Daisy: An efficient dense descriptor applied to wide-baseline stereo." Pattern Analysis and Machine Intelligence, IEEE Transactions


Common operations: DAISY

Num histograms



Num rings

num orientations (bins)




Params:

step, radius, num rings,
num histograms per ring,
orientations per histogram

Common operations: DAISY


Num histograms



Num rings



num Bag of Features Image Representation

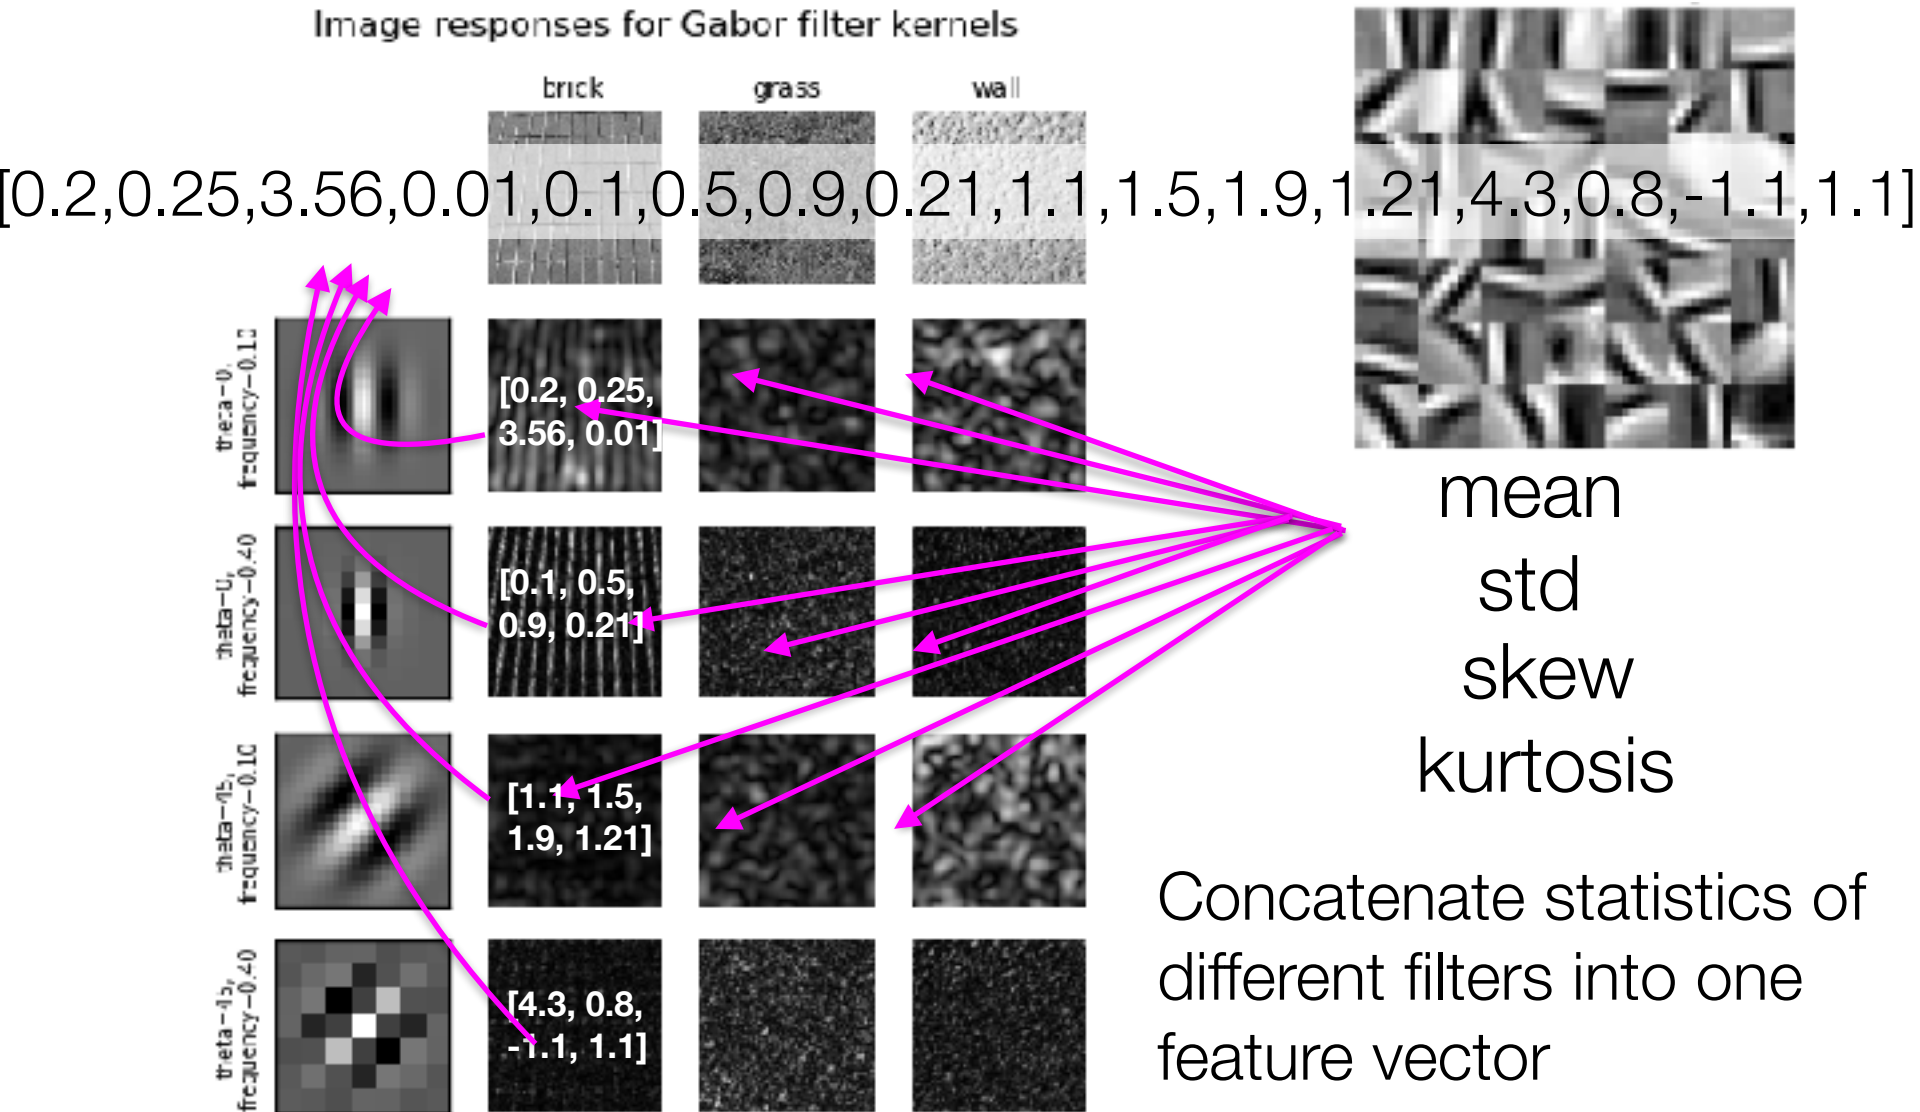


Params:

step, radius, num rings,
num histograms per ring,
orientations per histogram

Common operations: Gabor filter Banks (if time)

- common used for texture classification



More Image Processing

Gradients

DAISY

Gabor Filter Banks



Other Tutorials:

http://scikit-image.org/docs/dev/auto_examples/

For Next Lecture

- Work on your text datasets!
- **Next Time:** No Class, Project work day
- **Next Week:** In-Class Assignment One!!!