

Lecture Notes for **Machine Learning in Python**

Professor Eric Larson
Logistic Regression

Class Logistics and Agenda

- Logistics
 - A2 Due soon!
- Agenda
 - Town Hall
 - Logistic Regression
 - Solving
 - Programming
 - Finally some real python!

Town Hall



Matching versus Bag of Features

- Not a difference of vectors, but a percentage of matching points



- SURF, ORB, SIFT, DAISY

Solving Logistic Regression



Setting Up Binary Logistic Regression

- From flipped lecture:

$\hat{y} = \mathbf{w}^T \hat{\mathbf{x}}$

$$P(y^{(i)} | \mathbf{x}^{(i)}, \mathbf{w}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x}^{(i)})}$$
$$P(y^{(i)} = 0 | \mathbf{x}^{(i)}, \mathbf{w}) = 1 - \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x}^{(i)})}$$

$\text{MAX } L(\mathbf{w})$
 $\mathbf{w}^* = \underset{\mathbf{w}}{\text{ARGMAX}} L(\mathbf{w})$


$$L(\mathbf{w}) = \prod_i g(\mathbf{w}^T \mathbf{x}^{(i)})_{y=1} \cdot (1 - g(\mathbf{w}^T \mathbf{x}^{(i)}))_{y=0}$$

maximize!

where $g(\cdot)$ is a sigmoid

How do you optimize iteratively?

- **Objective Function:** the function we want to minimize or maximize
- **Update Formula:** what update “step” can we take to optimize the objective function
- **Parameters:** what are the parameters of the model that we can change to optimize the objective function



Binary Solution for Update Equation

- Video Supplement (also on canvas):
 - <https://www.youtube.com/watch?v=FGnoHdjFrJ8>
- General Procedure:
 - Simplify $L(\mathbf{w})$ with logarithm, $l(\mathbf{w})$

$$l(\mathbf{w}) = \sum_i \mathbf{y}^{(i)} \ln (g(\mathbf{w}^T \mathbf{x}^{(i)})) + (1 - \mathbf{y}^{(i)}) \ln (1 - g(\mathbf{w}^T \mathbf{x}^{(i)}))$$

- Take Gradient

$$\frac{\partial}{\partial w_j} l(\mathbf{w}) = - \sum_i (\mathbf{y}^{(i)} - g(\mathbf{w}^T \mathbf{x}^{(i)})) x_j^{(i)}$$

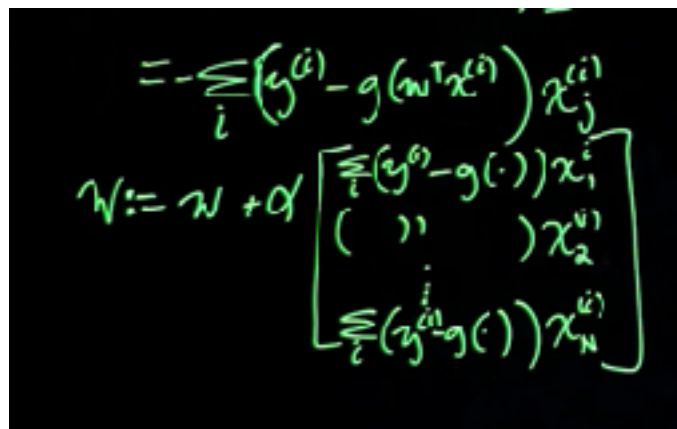
- Use gradient inside update equation for \mathbf{w}

Binary Solution for Update Equation

- Use gradient inside update equation for \mathbf{w}

$$\frac{\partial}{\partial w_j} l(\mathbf{w}) = - \sum_i (y^{(i)} - g(\mathbf{w}^T \mathbf{x}^{(i)})) x_j^{(i)}$$

$$\underbrace{w_j}_{\text{new value}} \leftarrow \underbrace{w_j}_{\text{old value}} + \underbrace{\eta \sum_{i=1}^M (y^{(i)} - g(\mathbf{w}^T \mathbf{x}^{(i)})) x_j^{(i)}}_{\text{gradient}}$$



Handwritten green text on a black background showing the gradient and update equation for weights:

$$= - \sum_i (y^{(i)} - g(\mathbf{w}^T \mathbf{x}^{(i)})) x_j^{(i)}$$
$$\mathbf{w} := \mathbf{w} + \alpha \begin{bmatrix} \sum_i (y^{(i)} - g(\cdot)) x_1^{(i)} \\ \vdots \\ \sum_i (y^{(i)} - g(\cdot)) x_n^{(i)} \end{bmatrix}$$

05. Logistic Regression.ipynb

Programming
Vectorization
Regularization
Multi-class extension



Other Tutorials:

<http://blog.yhat.com/posts/logistic-regression-python-rodeo.html>

http://scikit-learn.org/stable/auto_examples/linear_model/plot_iris_logistic.html

For Next Lecture

- **Next time:** More gradient based optimization techniques for logistic regression


Lecture Notes for **Machine Learning in Python**

Professor Eric Larson
Optimization Techniques for Logistic Regression

Class Logistics and Agenda

- Agenda
 - Finish Logistic Regression
 - More Numerical Optimization Techniques
 - Types of Optimization
 - Programming the Optimization
- **Whirlwind Lecture Alert**
 - Get an intuition, program it, maybe you don't follow every mathematical concept in lecture
 - But you know how to approach it outside lecture

What do weights and intercept define?



05. Logistic Regression.ipynb


“Finish”

Programming

Vectorization

Regularization

Multi-class extension



Other Tutorials:

<http://blog.yhat.com/posts/logistic-regression-python-rodeo.html>

http://scikit-learn.org/stable/auto_examples/linear_model/plot_iris_logistic.html

Demo Lecture

06. Optimization



Scratch Paper

Back Up Slides

Last time

$$p(y^{(i)} = 1 \mid x^{(i)}, w) = \frac{1}{1 + \exp(w^T x^{(i)})}$$


$$l(w) = \sum_i (y^{(i)} \ln[g(w^T x^{(i)})] + (1 - y^{(i)}) (\ln[1 - g(w^T x^{(i)})]))$$

$$\underbrace{w_j}_{\text{new value}} \leftarrow \underbrace{w_j}_{\text{old value}} + \underbrace{\eta \sum_{i=1}^M (y^{(i)} - g(x^{(i)})) x_j^{(i)}}_{\text{gradient}}$$

$$w \leftarrow w + \eta \sum_{i=1}^M (y^{(i)} - g(x^{(i)})) x^{(i)}$$

$$w \leftarrow w + \eta \left[\underbrace{\nabla l(w)_{\text{old}}}_{\text{old gradient}} - C \cdot 2w \right]$$

```
def _get_gradient(self, X, y):  
    # programming \sum_i (yi - g(xi)) xi  
    gradient = np.zeros(self.w_.shape) # set  
    for (xi, yi) in zip(X, y):  
        # the actual update inside of sum  
        gradi = (yi - self.predict_proba(xi,  
        # reshape to be column vector and add  
        gradient += gradi.reshape(self.w_.sh  
  
    return gradient/float(len(y))
```




Optimization: gradient descent

- What we know thus far:

$$\underbrace{w_j}_{\text{new value}} \leftarrow \underbrace{w_j}_{\text{old value}} + \eta \underbrace{\left[\left(\sum_{i=1}^M (y^{(i)} - g(x^{(i)})) x_j^{(i)} \right) - C \cdot 2w_j \right]}_{\nabla l(w)}$$

$$w \leftarrow w + \eta \nabla l(w)$$




Line Search: a better method

- Line search in direction of gradient:

$$\eta \leftarrow \arg \max_{\eta} \underbrace{\sum_{i=1}^M (y^{(i)} - \hat{y}^{(i)})^2 - C \cdot \sum_j w_j^2}_{\nabla l(w)}$$

$$w \leftarrow w + \eta \nabla l(w)$$

$$w \leftarrow w + \underbrace{\eta}_{\text{best step?}} \nabla l(w)$$



Revisiting the Gradient

- How much computation is required (for gradient)?

$$\sum_{i=1}^M (y^{(i)} - \hat{y}^{(i)})x^{(i)} - 2C \cdot w$$

M = number of instances
N = number of features

**Self Test: How many multiplies
per gradient calculation?**

- A. $M \cdot N + 1$ multiplications
- B. $(M+1) \cdot N$ multiplications
- C. $2N$ multiplications
- D. $2N - M$ multiplications

Stochastic Methods

- How much computation is required (for gradient)?

$$\sum_{i=1}^M (y^{(i)} - \hat{y}^{(i)})x^{(i)} - 2C \cdot w$$

Per iteration:


(M+1)*N multiplications
2M add/subtract

$$w \leftarrow w + \underbrace{\eta \left((y^{(i)} - \hat{y}^{(i)})x^{(i)} - 2C \cdot w \right)}_{\text{approx. gradient}},$$

i chosen at random

Per iteration:

N+1 multiplications
1 add/subtract



Gradient Descent (with line search)
Stochastic Gradient Descent
Hessian
Quasi-Newton Methods
Multi-processing




For Next Lecture

- **Next time:** SVMs via in class assignment
- **Next Next time:** Neural Networks



Can we do better than the gradient?

- Assume function is quadratic:



function of one variable:

$$w \leftarrow w - \underbrace{\left[\frac{\partial^2}{\partial w} l(w) \right]^{-1}}_{\text{inverse 2nd deriv}} \underbrace{\frac{\partial}{\partial w} l(w)}_{\text{derivative}}$$

will solve in one step!

what is the second order derivative
for a multivariate function?

$$\nabla^2 l(w) = \mathbf{H}[l(w)]$$

The Hessian

- Assume function is quadratic:

function of one variable:

$$\mathbf{H}[l(w)] = \begin{bmatrix} \frac{\partial^2}{\partial w_1} l(w) & \frac{\partial}{\partial w_1} \frac{\partial}{\partial w_2} l(w) & \dots & \frac{\partial}{\partial w_1} \frac{\partial}{\partial w_N} l(w) \\ \frac{\partial}{\partial w_2} \frac{\partial}{\partial w_1} l(w) & \frac{\partial^2}{\partial w_2} l(w) & \dots & \frac{\partial}{\partial w_2} \frac{\partial}{\partial w_N} l(w) \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial}{\partial w_N} \frac{\partial}{\partial w_1} l(w) & \frac{\partial}{\partial w_N} \frac{\partial}{\partial w_2} l(w) & \dots & \frac{\partial^2}{\partial w_N} l(w) \end{bmatrix}$$




$$\nabla^2 l(w) = \mathbf{H}[l(w)]$$

The Newton Update Method

- Assume function is quadratic (in high dimensions):


$$w \leftarrow w - \underbrace{\left[\frac{\partial^2}{\partial w} l(w) \right]^{-1}}_{\text{inverse 2nd deriv}} \underbrace{\frac{\partial}{\partial w} l(w)}_{\text{derivative}}$$

$$w \leftarrow w + \eta \cdot \underbrace{\mathbf{H}[l(w)]^{-1}}_{\text{inverse Hessian}} \cdot \underbrace{\nabla l(w)}_{\text{gradient}}$$



Is. Newton

I do not know what I may appear to the world, but to myself I seem to have been only like a boy playing on the sea-shore, and diverting myself in now and then finding a smoother pebble or a prettier shell than ordinary, whilst the great ocean of truth lay all undiscovered before me.



I do not know what I may appear to the world, but to myself I seem to have been only like a boy playing on the sea-shore, and diverting myself in now and then finding a smoother pebble or a prettier shell than ordinary, whilst the great ocean of truth lay all undiscovered before me.

Is. Newton

$$\frac{\partial}{\partial w_j} l(w) = \sum_i (y^{(i)} - g(x^{(i)})) x_j^{(i)}$$

↓ PLUG IN

$$H[k,j] = \frac{\partial}{\partial w_k} \frac{\partial}{\partial w_j} l = \frac{\partial}{\partial w_k} \left(\sum_i (y^{(i)} - g(x^{(i)})) x_j^{(i)} \right)$$

$$= \sum_i \frac{\partial}{\partial w_k} y^{(i)} x_j^{(i)} - \frac{\partial}{\partial w_k} g(x^{(i)}) x_j^{(i)}$$

← LEFT OVER TERM

$$= - \sum_i \frac{\partial}{\partial w_k} g(x^{(i)}) x_j^{(i)}$$

Already know $\frac{\partial}{\partial w_k} g(w^T x^{(i)})$ side calculation

$$= g(x^{(i)}) (1 - g(x^{(i)})) \frac{\partial}{\partial w_k} (w^T x^{(i)})$$

$$= g(x^{(i)}) (1 - g(x^{(i)})) x_k^{(i)}$$

← PLUG IN

$$\therefore = - \sum_i g(x^{(i)}) (1 - g(x^{(i)})) x_k^{(i)} x_j^{(i)}$$

← THAT'S THE HESSIAN!

$$H(k,j) =$$

This is a valid equation for the Hessian, but we want to represent it using linear algebra

$$= - \sum_i g(x^{(i)}) (1 - g(x^{(i)})) x_k^{(i)} x_j^{(i)}$$

↑ SIGMOID

3D LINEAR ALG.

The Hessian for Logistic Regression

- The hessian is easy to calculate from the gradient for logistic regression

$$w \leftarrow w + \eta \cdot \underbrace{\mathbf{H}[l(w)]^{-1}}_{\text{inverse Hessian}} \cdot \underbrace{\nabla l(w)}_{\text{gradient}}$$

$$\mathbf{H}_{j,k}[l(w)] = - \sum_{i=1}^M g(x^{(i)})(1 - g(x^{(i)})) x_k^{(i)} x_j^{(i)}$$

$$\underbrace{\sum_{i=1}^M (y^{(i)} - \hat{y}^{(i)}) x_j^{(i)}}_{\text{gradient}}$$

$$\mathbf{H}[l(w)] = X^T \cdot \text{diag}[g(x^{(i)})(1 - g(x^{(i)}))] \cdot X$$

$$X * y_{diff}$$

$$w \leftarrow w + \eta [X^T \cdot \text{diag}[g(x^{(i)})(1 - g(x^{(i)}))] \cdot X]^{-1} \cdot X * y_{diff}$$

Newton's method



Problems with Newton's Method

- **Quadratic** isn't always a great assumption:
 - highly dependent on starting point
 - jumps can get **really random!**
 - near saddle points, inverse hessian **unstable**
 - hessian **not** always **invertible**...
 - or invertible with correct numerical precision

The solution: quasi Newton methods

- In general:
 - **approximate** the **Hessian** with something numerically sound and efficiently invertible
 - **back off to gradient** descent when the approximate hessian is **not stable**
 - use **momentum** to update approximate hessian
- **A popular approach:** use Broyden-Fletcher-Goldfarb-Shanno (BFGS)
 - which you can look up if you are interested ...

https://en.wikipedia.org/wiki/Broyden-Fletcher-Goldfarb-Shanno_algorithm

BFGS

$$\mathbf{H}_0 = \mathbf{I} \quad \text{init}$$

$$p_k = -\mathbf{H}_k^{-1} \nabla l(w_k) \quad \text{get update direction}$$

$$\text{find next } w \quad w_{k+1} \leftarrow w_k + \eta \cdot p_k$$

$$\text{get scaled direction} \quad s_k = \eta \cdot p_k$$

$$v_k = \nabla l(w_{k+1}) - \nabla l(w_k) \quad \text{approx gradient change}$$

$$\mathbf{H}_{k+1} = \mathbf{H}_k + \underbrace{\frac{v_k v_k^T}{v_k^T s_k}}_{\text{approx. Hessian}} - \underbrace{\frac{\mathbf{H}_k s_k s_k^T \mathbf{H}_k}{s_k^T \mathbf{H}_k s_k}}_{\text{momentum}} \quad \text{update Hessian and inverse Hessian approx}$$

$$\mathbf{H}_{k+1}^{-1} = \mathbf{H}_k^{-1} + \frac{(s_k^T v_k + \mathbf{H}_k^{-1})(s_k s_k^T)}{(s_k^T v_k)^2} - \frac{\mathbf{H}_k^{-1} v_k s_k^T + s_k v_k^T \mathbf{H}_k^{-1}}{s_k^T v_k}$$

$$k = k + 1 \quad \text{increment } k \text{ and repeat}$$

invertibility of \mathbf{H} well defined / only matrix operations

BFGS (if time)
parallelization

