

---

# Lecture Notes for Machine Learning in Python

---

Professor Eric Larson  
Week Nine A

# Class Logistics and Agenda

---

- Grades Coming Soon, but slowly
- **Next week:** project work week
- Agenda:
  - More advanced Neural Network Architectures
  - Ensemble methods
- **Next Time:** in-class assignment

---

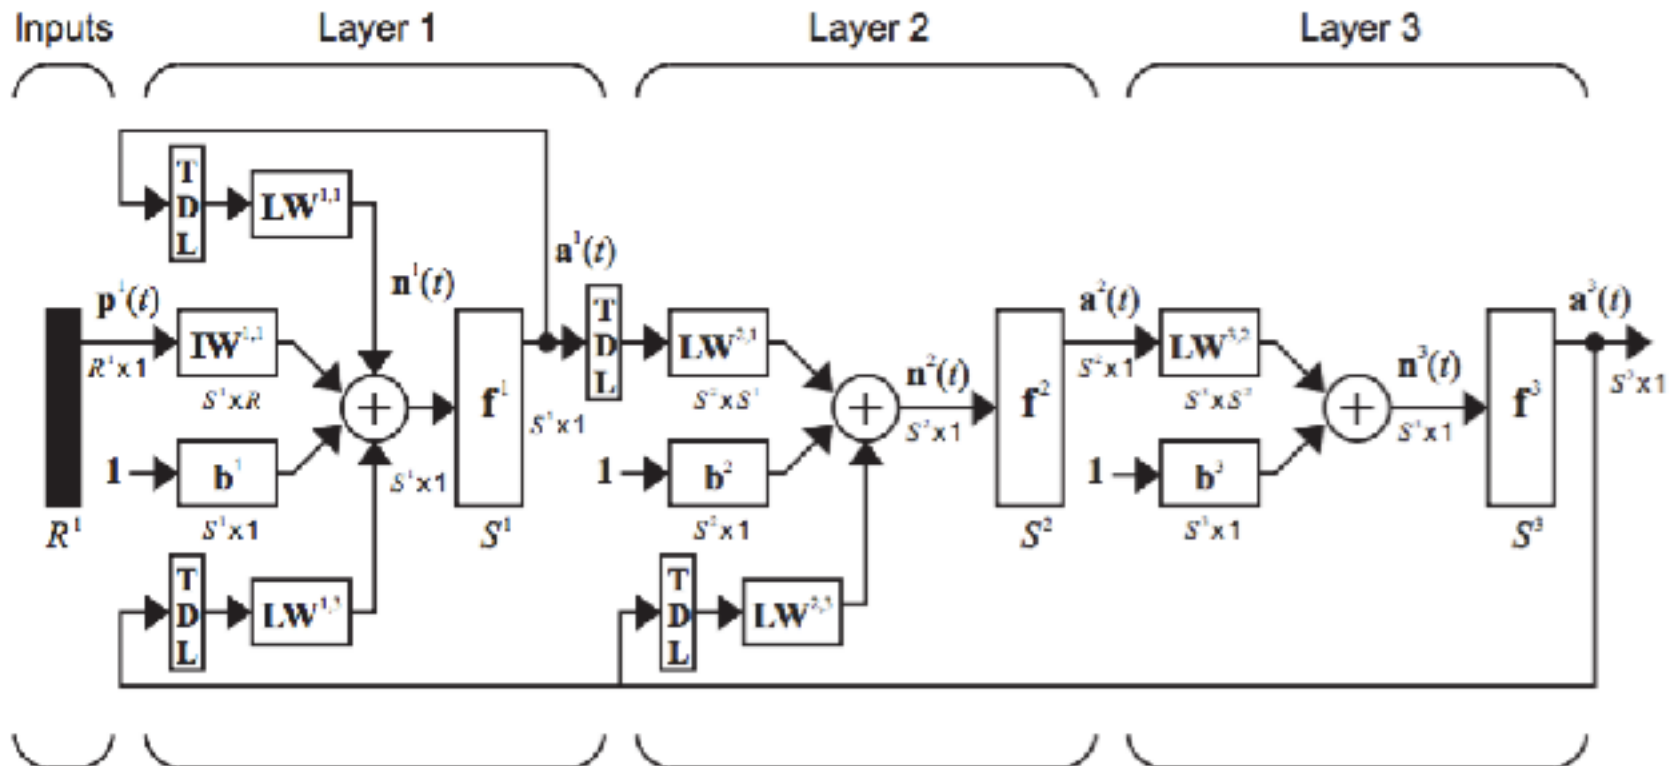
# Even More Advanced Architectures

---



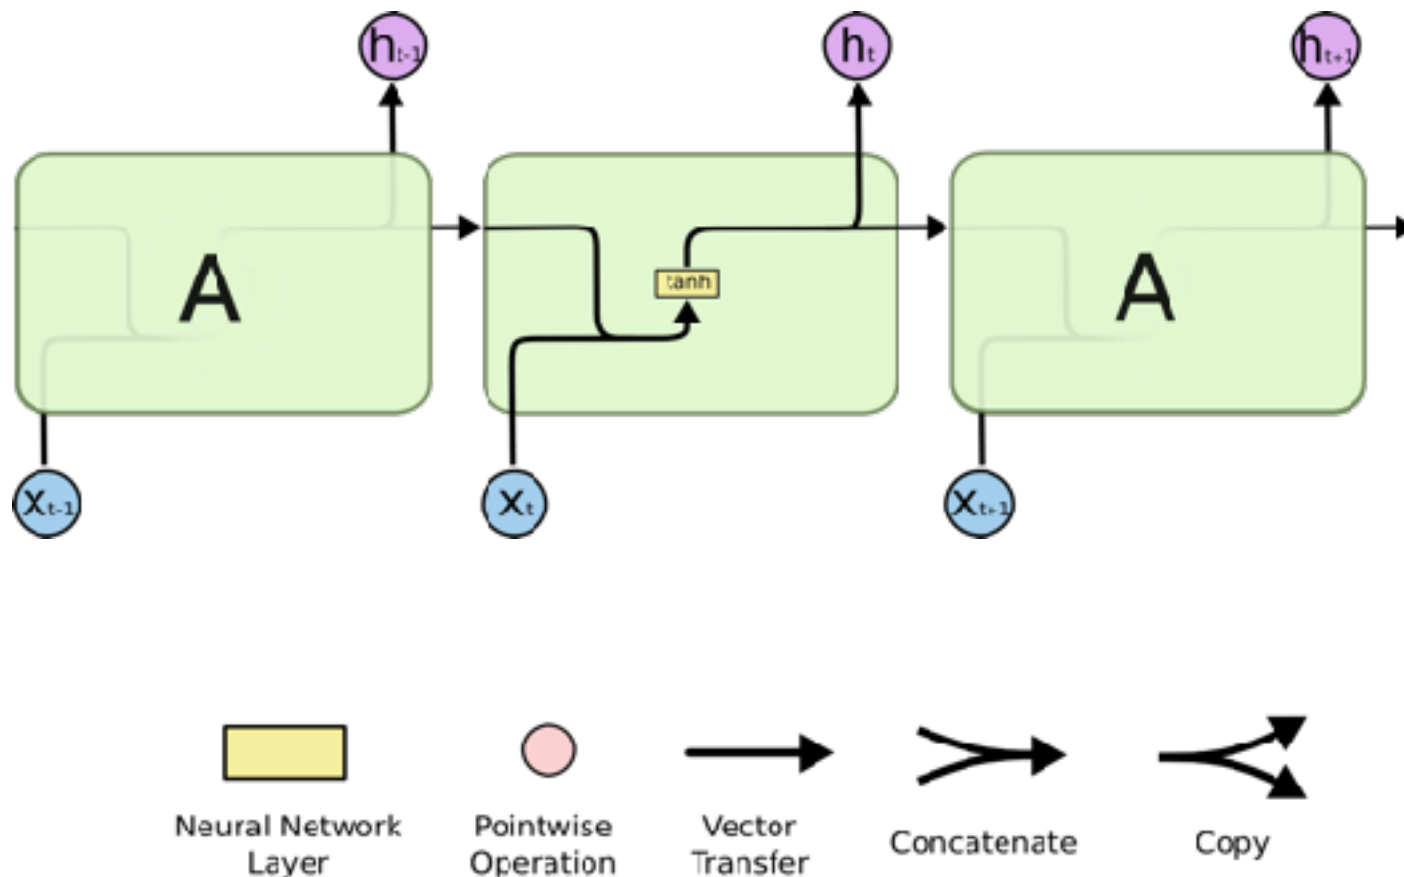
# More Advanced Architectures

- Dynamic Networks (recurrent networks)
  - can use current and previous inputs, in time
  - still popular, but ultimately extremely hard to train
  - **highly successful variant:** long short term memory, **LSTM**



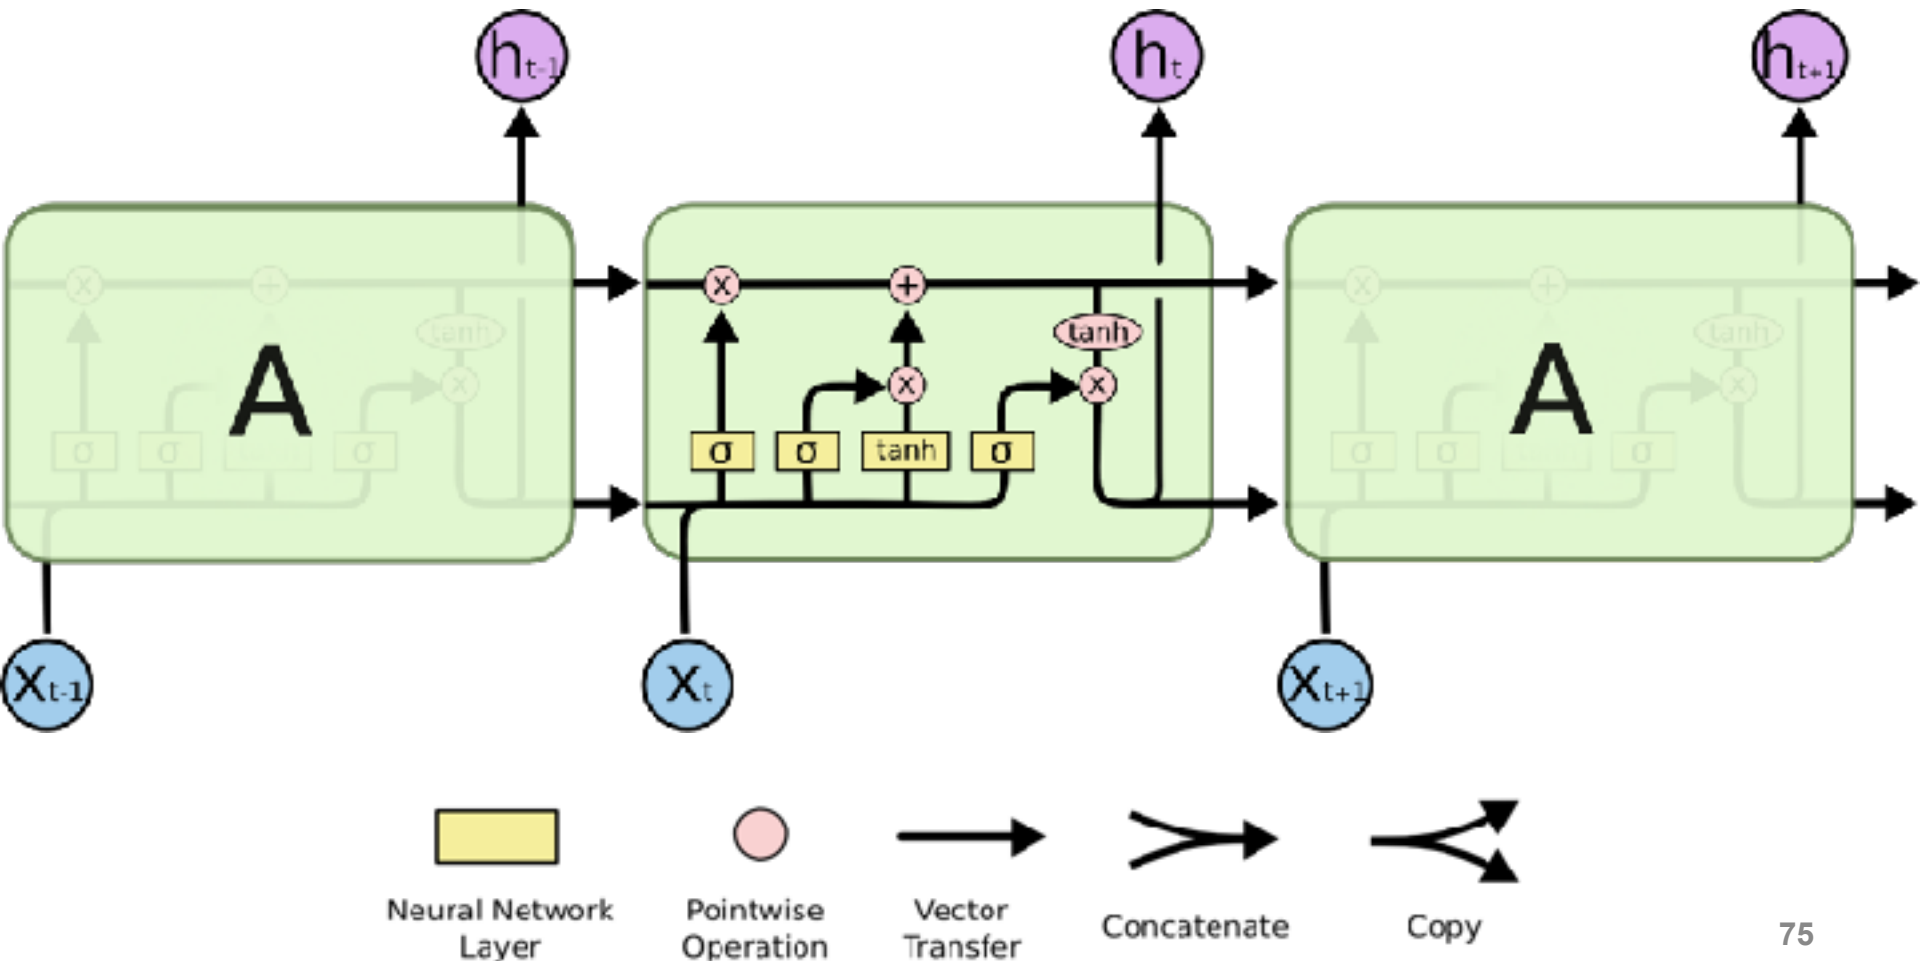
# More Advanced Architectures

- **LSTM key idea:** limit how past data can affect output




# More Advanced Architectures

- **LSTM key idea:** limit how past data can affect output




# More Advanced Architectures


- **LSTM key idea:** limit how past data can affect output



potentially **forget past** inputs  
via “gate”  $\sigma$

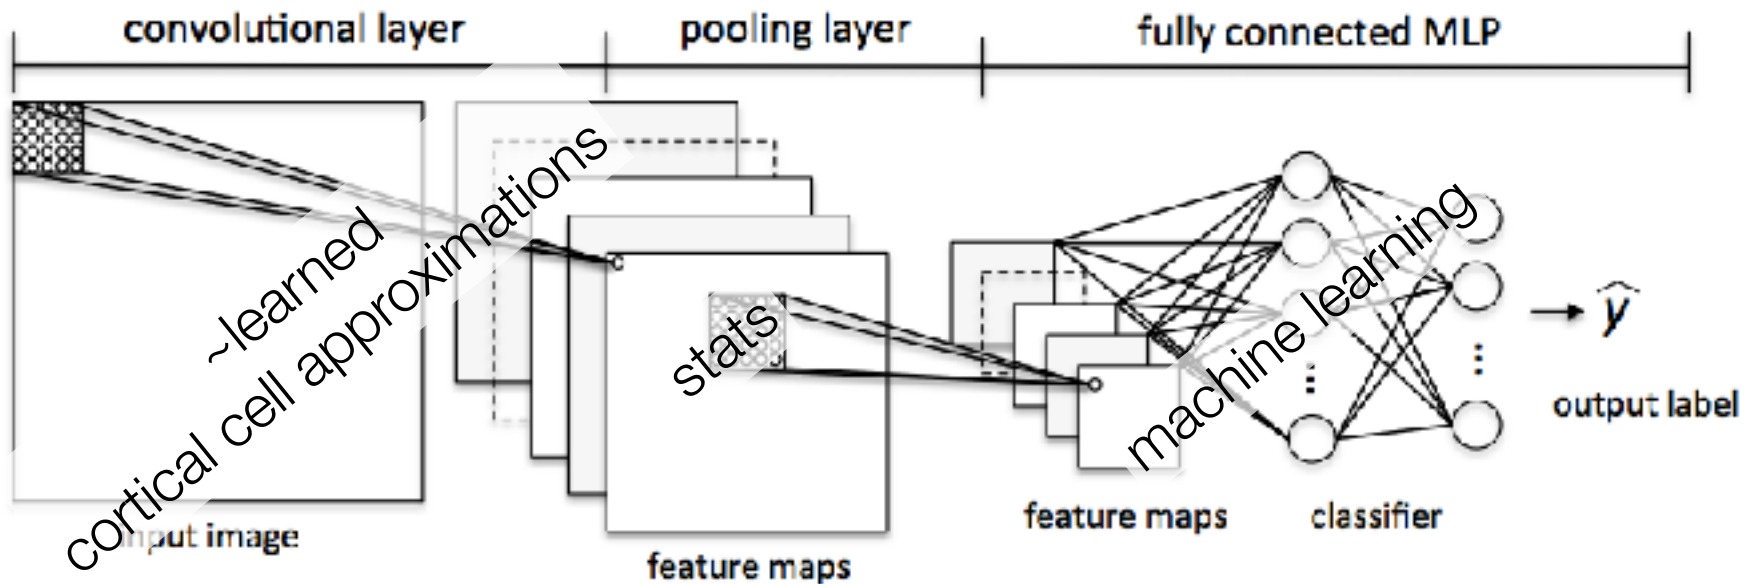


we will return to this architecture later, for now:  
**put it in long term memory** 😂



# More Advanced Architectures

- Convolutional Neural Networks
  - image processing operations



we will return to this architecture later, for now:  
**put it in long term memory**



# Problems with these Advanced Architectures

---

- These architectures have been around for 30 years
- And solved some amazingly hard problems
- but they had **big training problems** that back propagation could not solve readily:
  - unstable gradients (vanishing/exploding)
  - **extremely** non-convex space
    - more layers==many more local optima to get stuck
  - sometimes **gradient optimization is too computational** for weight updates
    - might need better optimization strategy than SGD
- The solution to these problems came from having large amounts of training data, better setup of the optimization
  - eventually was termed **deep learning**

# End of Neural Networks Introduction

---

- Briefly step away from Neural Networks
- **Next time: In class assignment:**
  - evaluation and cross validation
- **Next Next time:** Ensembles

these will relate  
back to neural networks

## More help on neural networks:

Sebastian Raschka

<https://github.com/rasbt/python-machine-learning-book/blob/master/code/ch12/ch12.ipynb>

Martin Hagan

[https://www.google.com/url?sa=t&rct=j&q=&esrc=s&source=web&cd=1&cad=rja&uact=8&ved=0ahUKEwioprwn27fPAhWMx4MKHYbwDIwQFggeMAA&url=http%3A%2F%2Fhagan.okstate.edu%2FNNDesign.pdf&usg=AFQjCNG5YbM4xSMm6K5HNsG-4Q8TvOu\\_Lw&sig2=bgT3k-5ZDDTPZ07Qu8Oreg](https://www.google.com/url?sa=t&rct=j&q=&esrc=s&source=web&cd=1&cad=rja&uact=8&ved=0ahUKEwioprwn27fPAhWMx4MKHYbwDIwQFggeMAA&url=http%3A%2F%2Fhagan.okstate.edu%2FNNDesign.pdf&usg=AFQjCNG5YbM4xSMm6K5HNsG-4Q8TvOu_Lw&sig2=bgT3k-5ZDDTPZ07Qu8Oreg)


Michael Nielsen

<http://neuralnetworksanddeeplearning.com>

---

## Classification: Ensemble Methods

---

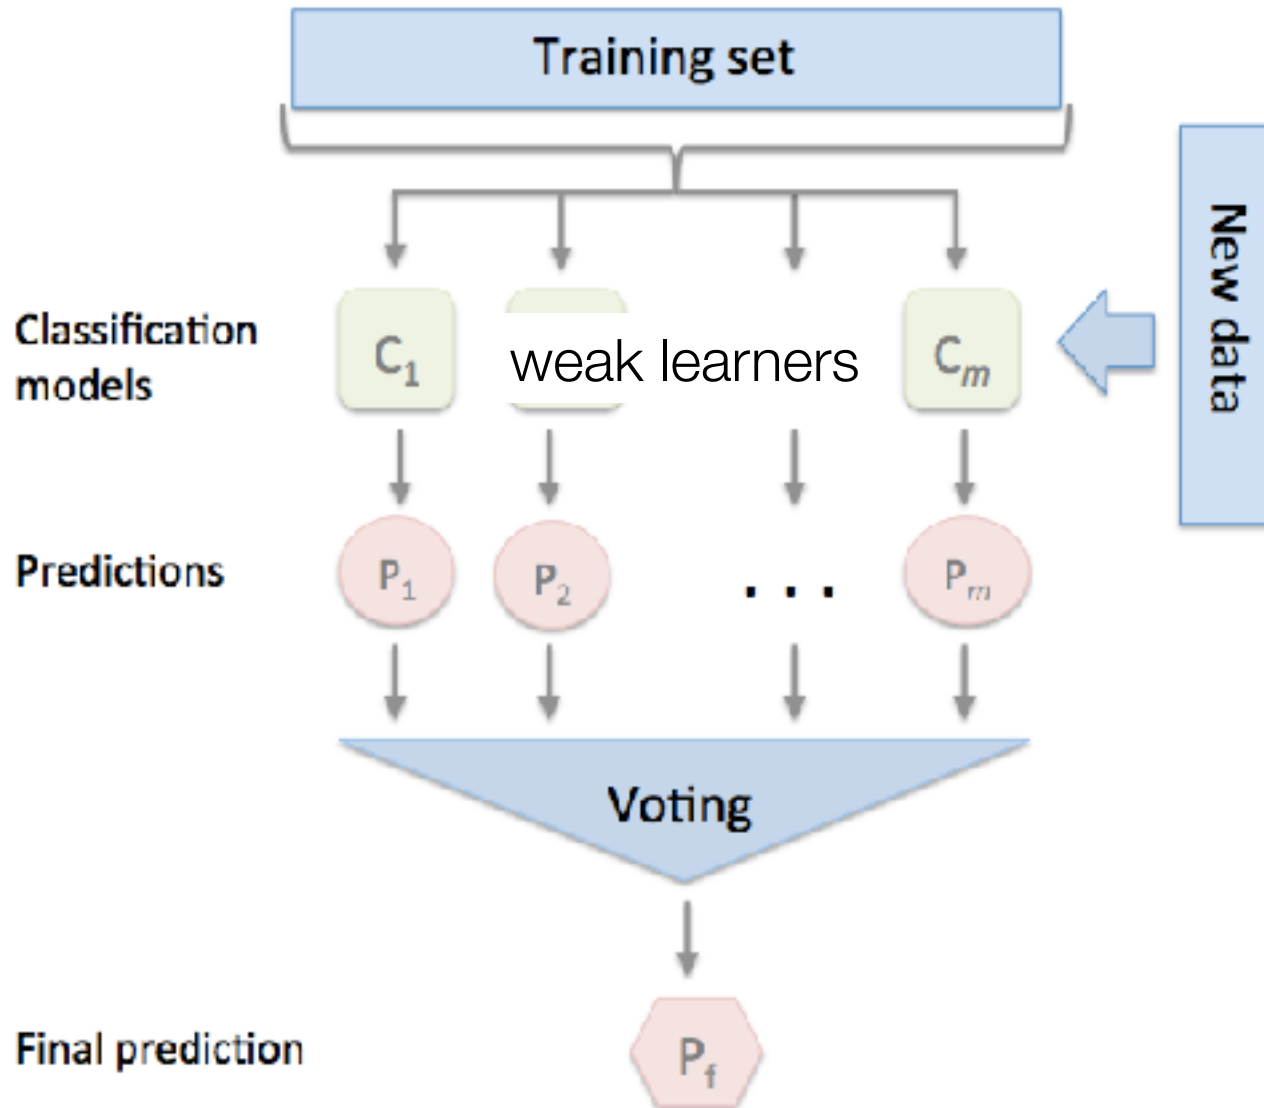


# Ensemble Methods

---

- Construct a set of classifiers from the training data
- Predict class label of previously unseen records by aggregating predictions made by multiple classifiers
- Could be multiple Neural Networks

# General Idea



# Weak and Strong

---


- **Weak learner:** a learner with better than chance accuracy (error  $< 50\%$  for two class problem)
  - *i.e.*, classifier has high bias
  - like logistic regression
- **Strong learner:** arbitrarily small error rate
  - like MLP or kernel SVM
- Need to Ensemble many weak learners

# Why does it work?

- Suppose there are 25 base classifiers
  - Each classifier has error rate,  $\varepsilon = 0.35$
  - Assume classifiers are independent, so they make errors on different samples from dataset
  - Probability that the ensemble classifier makes a wrong prediction:

$$\sum_{i=13}^{25} \binom{25}{i} \varepsilon^i (1 - \varepsilon)^{25-i} = 0.06$$

- But in practice, our classifiers are correlated, so it **does not work this well**



# Why does it work?

---

- How much does this horse weigh?
  - Average of the guesses from many people is close to the true value
  - Average of *many* people is better than an expert's guess

## **Self Test:**

- A. 250 lbs
- B. 750 lbs
- C. 1200 lbs
- D. 5000 lbs





# Ensembles of Different Classifiers

---

- Step one: train  $m$  classifiers from dataset,  $C_m(\mathbf{x})$
- Step two: combine outputs
  - majority vote:

$$\arg \max_i \sum_j w_j [C_j(\mathbf{x})=i]$$

trust in classifier

classifier selected i

- majority probabilistic vote:

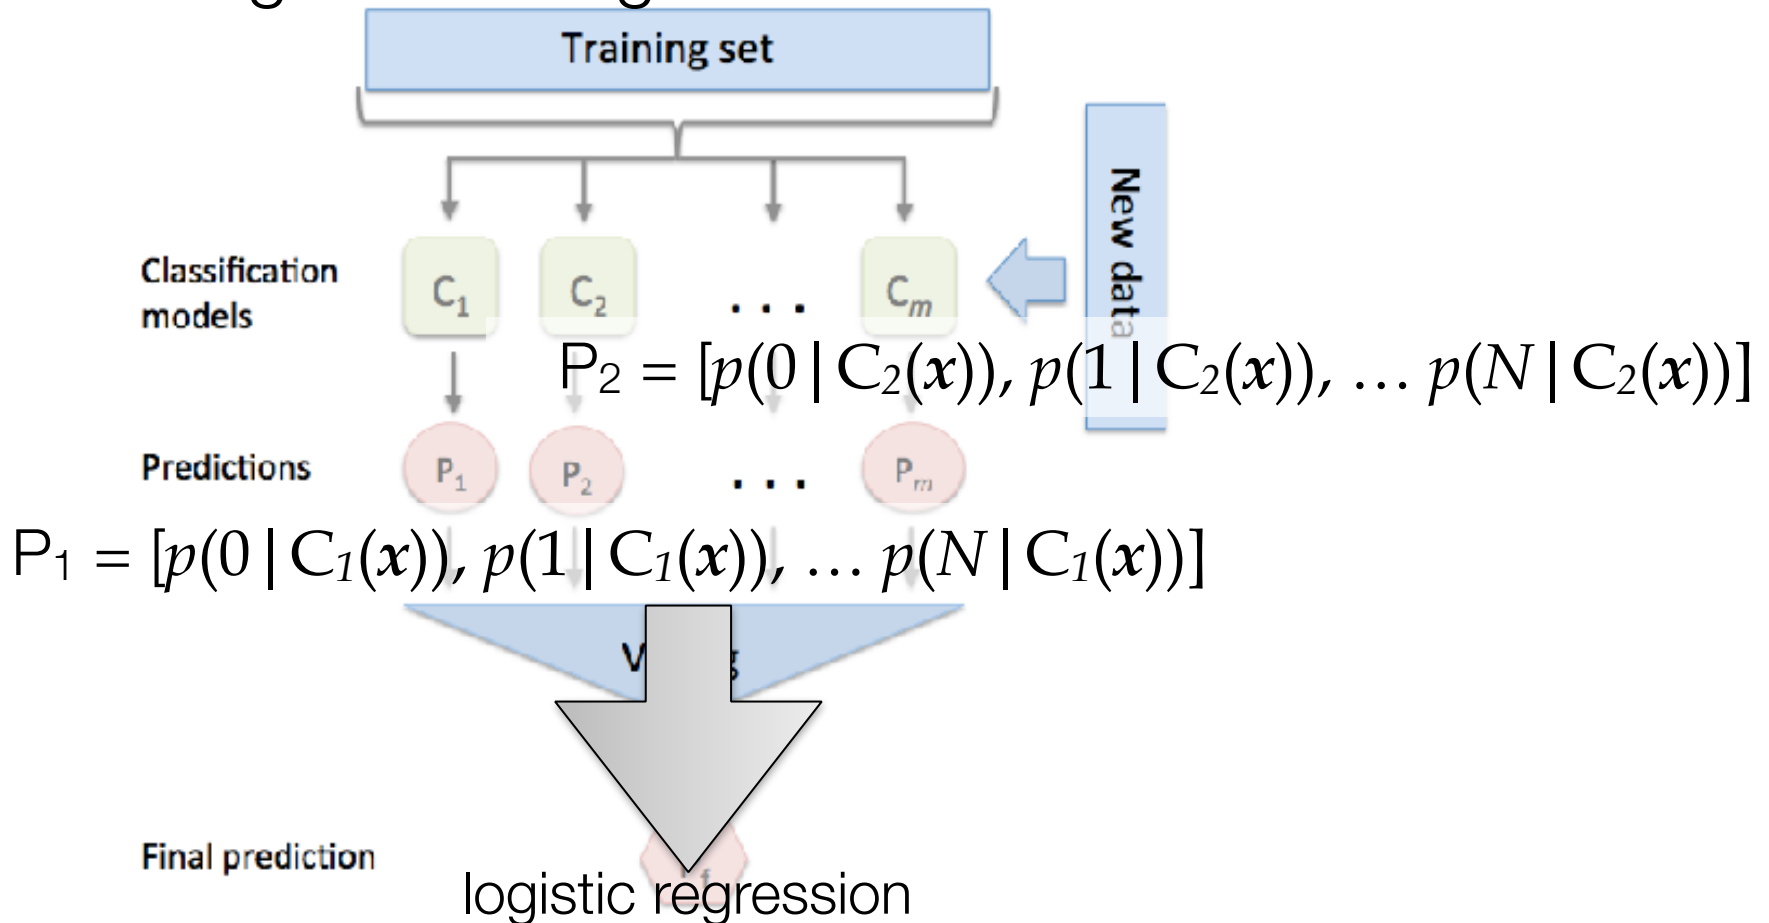
$$\arg \max_i \sum_j w_j p(i | C_j(\mathbf{x}))$$

trust in classifier

predict\_proba i

# Examples of Ensemble Methods

- Stacking/Cascading



# Examples of Ensemble Methods

---

- Training set sampling methods:
  - Bagging
  - Boosting

# Bagging

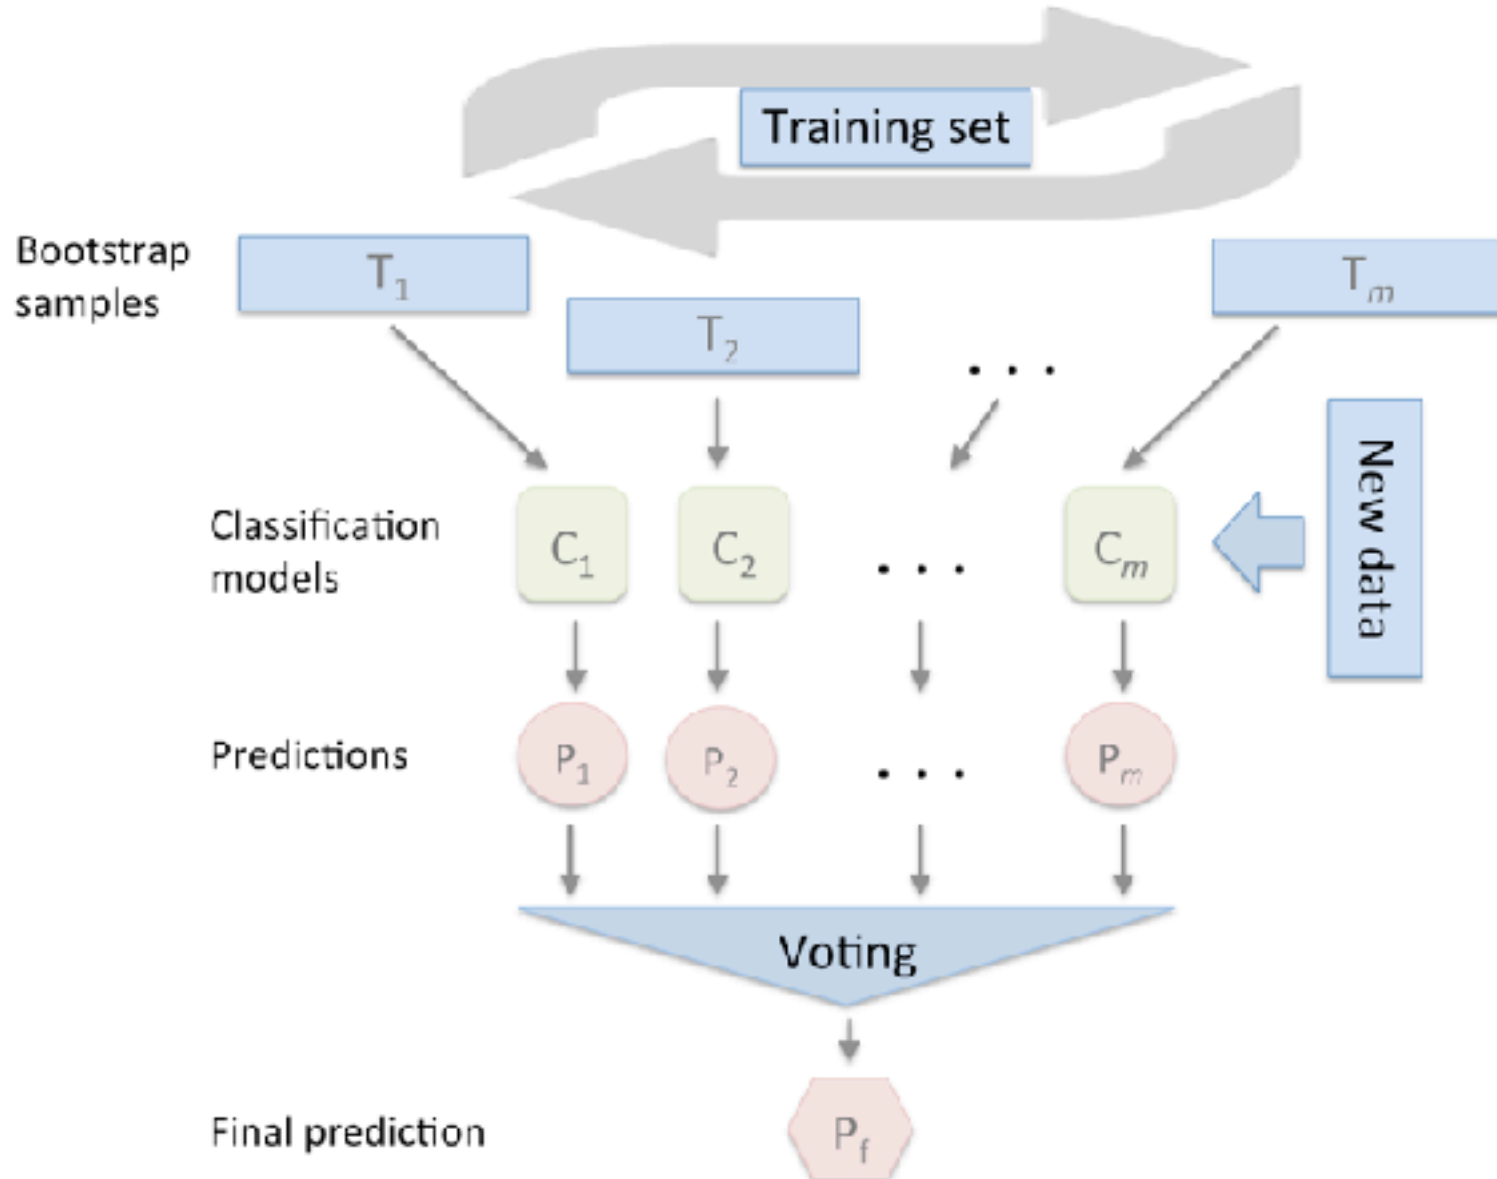
---

- Sampling with replacement

<b>Original Data</b>	1	2	3	4	5	6	7	8	9	10
<b>Bagging (Round 1)</b>	7	8	10	8	2	5	10	10	5	9
<b>Bagging (Round 2)</b>	1	4	9	1	2	3	2	7	3	2
<b>Bagging (Round 3)</b>	1	8	5	10	5	5	9	6	3	7


- Build classifiers from subsamples of data
  - could be smart or just completely random (uniform)
  - could sample from instances (rows) or from features (columns)
- Combine the resulting classifiers as before
  - majority vote
  - argmax probability

# Bagging



# The most famous bagging classifier

- Random Forests
  - Select random subset of samples
  - Select random subset of the features
  - build a tree
  - build many trees
  - actually a whole forest of trees




# Random Forest

---

- Random Forests
  - Decision trees are built
  - But at each stage, a random subset of the features is selected (random subspace)
    - if “ $f$ ” features, look at “ $np.sqrt(f)$ ” features at each iteration
  - Generalization built in: Out-of-bag
  - Variable importance:
    - random feature permutation
    - look at out-of-bag samples
    - randomly permute the values of  $n^{\text{th}}$  feature
    - see how performance degrades

# Random Forest



- One can use the training data to get an error estimate ("out of bag error" or OOBE)
- Validate each tree on complement of training data



# Features of Random Forests:

---

- produce high accuracy on many real world datasets
- run efficiently on large databases (each tree is an easy prediction, easily extensible to map reduce)
- can handle thousands of input variables without variable deletion
- give estimates of what variables are important in the classification
- generate an internal unbiased estimate of the generalization error as the forest building progresses
- have effective method for estimating missing data
- have methods for balancing error in class population for unbalanced data sets

# Boosting

---

- An iterative procedure to adaptively change distribution of training data by focusing more on previously misclassified records
  - Initially, all  $N$  records are assigned equal weights
  - Unlike bagging, weights may change at the end of boosting round
  - Samples with a higher weight are more likely to be chosen

# Boosting

---

- Records that are wrongly classified will have their weights increased
- Records that are classified correctly will have their weights decreased

Original Data	1	2	3	4	5	6	7	8	9	10
Boosting (Round 1)	7	3	2	8	7	9	4	10	6	3
Boosting (Round 2)	5	4	9	4	2	5	1	7	4	2
Boosting (Round 3)	4	4	8	10	4	5	4	6	3	4

- Example 4 is hard to classify
- Its weight is increased, therefore it is more likely to be chosen again in subsequent rounds

# Example: AdaBoost


- Base classifiers:  $C_1, C_2, \dots, C_T$
- Weighted Error rate of  $C_i$  is:

$$\epsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \delta(C_i(x_j) \neq y_j)$$

$j^{th}$  instance weight      indicator

- Importance of a classifier:

$$\alpha_i = \frac{1}{2} \ln \left( \frac{1 - \epsilon_i}{\epsilon_i} \right)$$



# Example: AdaBoost

---

- Weight update:

$$w_j \leftarrow w_j \times \begin{cases} 1 & \text{if } C_i(x_j) = y_i \\ (1 - \epsilon_i) / \epsilon_i & \text{if } C_i(x_j) \neq y_i \end{cases}$$

Decrease weight

Increase weight

- Rescale weights to sum to one
- Classification:

$$C^*(x) = \arg \max_y \sum_{j=1}^T \alpha_j \delta(C_j(x) = y)$$

because we take arg max,  
the 1/2 constant in  $\alpha$  is not needed

# Illustrating AdaBoost

- Original data (sorted on x):

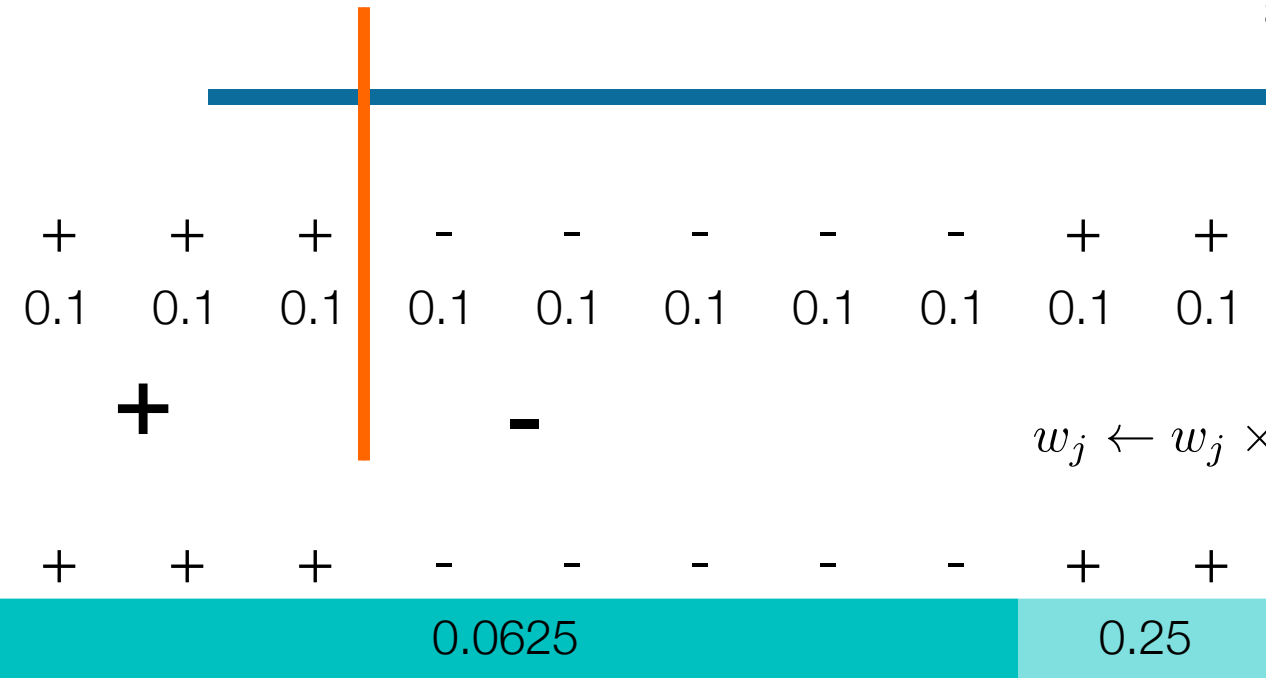
+	+	+	-	-	-	-	-	+	+
0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1
+		+		-	-		-		+
0.1		0.1		0.1	0.1		0.1		0.1
+			-						

- We have 10 data points, so each data point gets initial weight  $1/10$ .
- Suppose we sample *six* points
- Then train a “decision stump” classifier
- Which makes two errors with weight 0.1

$$\epsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \delta (C_i(x_j) \neq y_j)$$


$$\alpha_i = \frac{1}{2} \ln \left( \frac{1 - \epsilon_i}{\epsilon_i} \right)$$

$$w_j \leftarrow w_j \times \begin{cases} 1 & \text{if } C_i(x_j) = y_i \\ (1 - \epsilon_i) / \epsilon_i & \text{if } C_i(x_j) \neq y_i \end{cases}$$



- Which makes two errors with weight 0.1
- So  $\epsilon = 2 \times 0.1 = 0.2$ ,  $\alpha = \ln[(1 - 0.2) / 0.2] = \ln 4 \sim 1.38$
- So weights of incorrect answers get multiplied by 4
- Then weights are rescaled to sum to one

# Illustrating AdaBoost



- Sample six new samples, train stump
- Resulting in 3 errors with weights 0.0625
- So  $\epsilon = 3 \times 0.0625 = 0.1875$ ,
- $\alpha = \ln(1 - 0.1875) / 0.1875 = \ln 4.33 \sim 1.47$
- Update weights

$$w_j \leftarrow w_j \times \begin{cases} 1 & \text{if } C_i(x_j) = y_i \\ (1 - \epsilon_i) / \epsilon_i & \text{if } C_i(x_j) \neq y_i \end{cases}$$

$$\epsilon_i = \frac{1}{N} \sum_{j=1}^N w_j \delta(C_i(x_j) \neq y_j)$$

$$\alpha_i = \frac{1}{2} \ln \left( \frac{1 - \epsilon_i}{\epsilon_i} \right)$$



# Illustrating AdaBoost

---

- New weights are:

+	+	+	-	-	-	-	-	+	+
0.17			0.039				0.15		

- Chosen samples, round three

A	+	+	+	B		-		C	+	+	D
0.17						0.039		0.15			

- **Self test:** where is my new decision stump?

# Self test: Illustrating AdaBoost


---

- New weights are:

+	+	+	-	-	-	-	-	+	+
0.17			0.039				0.15		

- Chosen samples, round three

+	+	+			-		+	+
0.17			0.039				0.15	
							+	-



- So  $\varepsilon = 5 \times 0.039 = 0.195$ ,
- $\alpha = \ln(1 - 0.195) / 0.195 = \ln 4.13 \sim 1.42$

# Illustrating AdaBoost

---

- Combined classifiers:

• $C_1, \alpha=1.38$	+	+	+	-	-	-	-	-	-	-
• $C_2, \alpha=1.47$	-	-	-	-	-	-	-	-	+	+
• $C_3, \alpha=1.42$	+	+	+	+	+	+	+	+	+	+
• $C^*$	+	+	+	-	-	-	-	-	+	+

$$C^*(x) = \arg \max_y \sum_{j=1}^T \alpha_j \delta(C_j(x) = y)$$

# A final thought on boosting and bagging

---

- Boosting is what won the Netflix prize
- But was never implemented
  - “...additional accuracy gains that we measured did not seem to justify the engineering effort to bring them into a production environment.”

# Watch videos before class!

---

- Next time is an in-class-assignment
  - cross validation!