

---

# Lecture Notes for Machine Learning in Python

---

Professor Eric Larson  
Evaluation and Cross Validation, Video Lecture

---

# Model Evaluation Best Practices

---

---

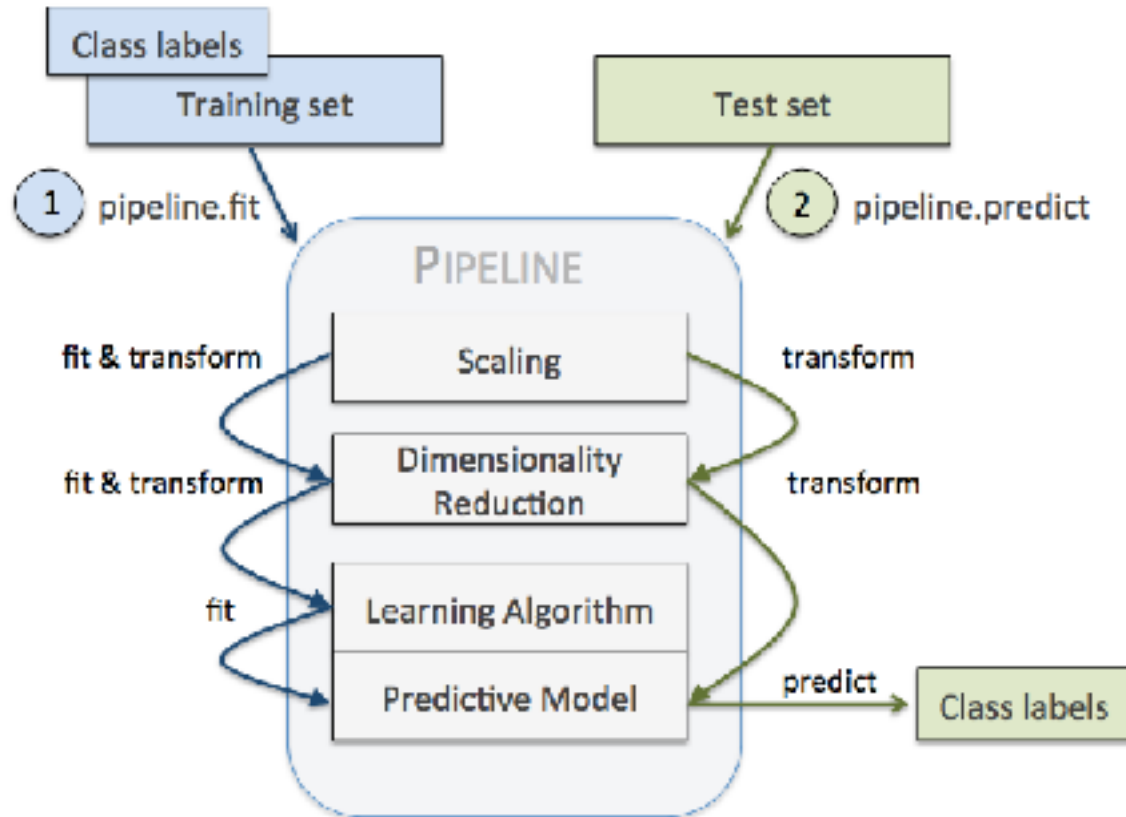
# Model Evaluation

---

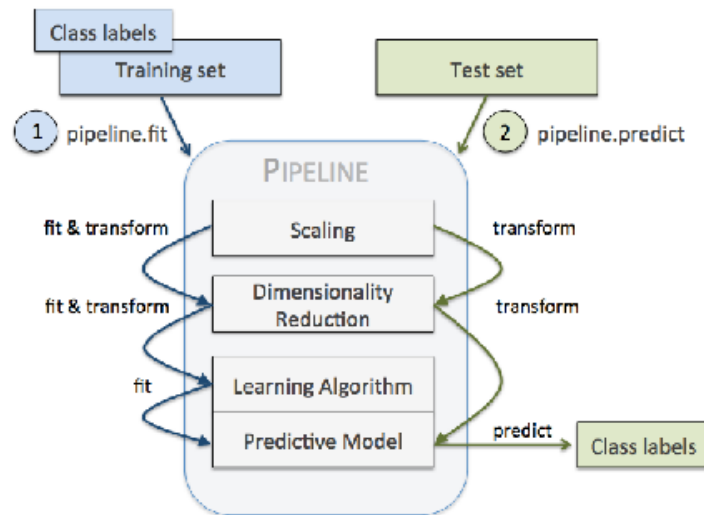
- How reliable are our estimates of performance?
- Performance of a model may depend on other factors besides the learning algorithm:
  - Class distribution
  - Cost of misclassification
  - Size of training and test sets

# Best Practice: Setup Pipelines

- Combine pre- and post-processing into stages
- Excellent way to prevent “data snooping”
  - guarantees separation of testing and training sets



# Best Practice: Setup Pipelines



```
pipe_lr = Pipeline([('scl', StandardScaler()),  
                    ('pca', PCA(n_components=2)),  
                    ('clf', LogisticRegression(random_state=1))])  
  
pipe_lr.fit(X_train, y_train)  
print('Test Accuracy: %.3f' % pipe_lr.score(X_test, y_test))  
y_pred = pipe_lr.predict(X_test)
```

**Test Accuracy: 0.947**

# Best Practice: Use Pipelines in Validation Loops

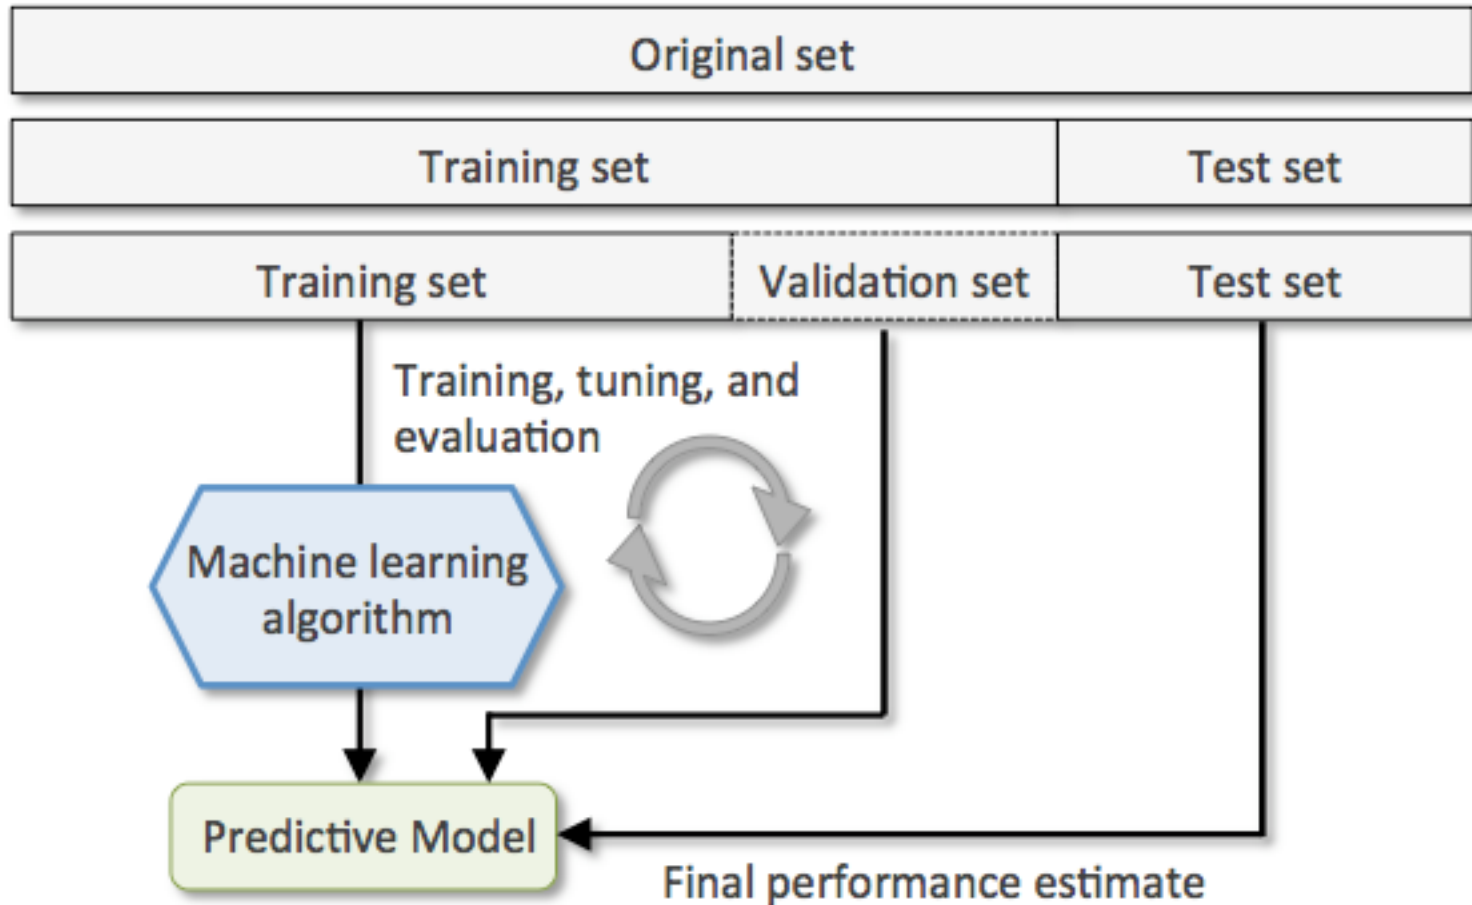
---

use testing set, and *never, never, never* let the model see it, Many different strategies:

- Holdout
  - Reserve  $x\%$  for training and  $(1-x)\%$  for testing
- Random subsampling
  - Repeated holdout, with replacement
- Cross validation
  - Partition data into  $k$  disjoint subsets
  - $k$ -fold: train on  $k-1$  partitions, test on the remaining one
  - Leave-one-out:  $k=M$
- Stratified Cross Validation
  - Select samples, keeping overall class distribution same for each fold

# Validation Loop Strategies

## Holdout



# Validation Loop Strategies

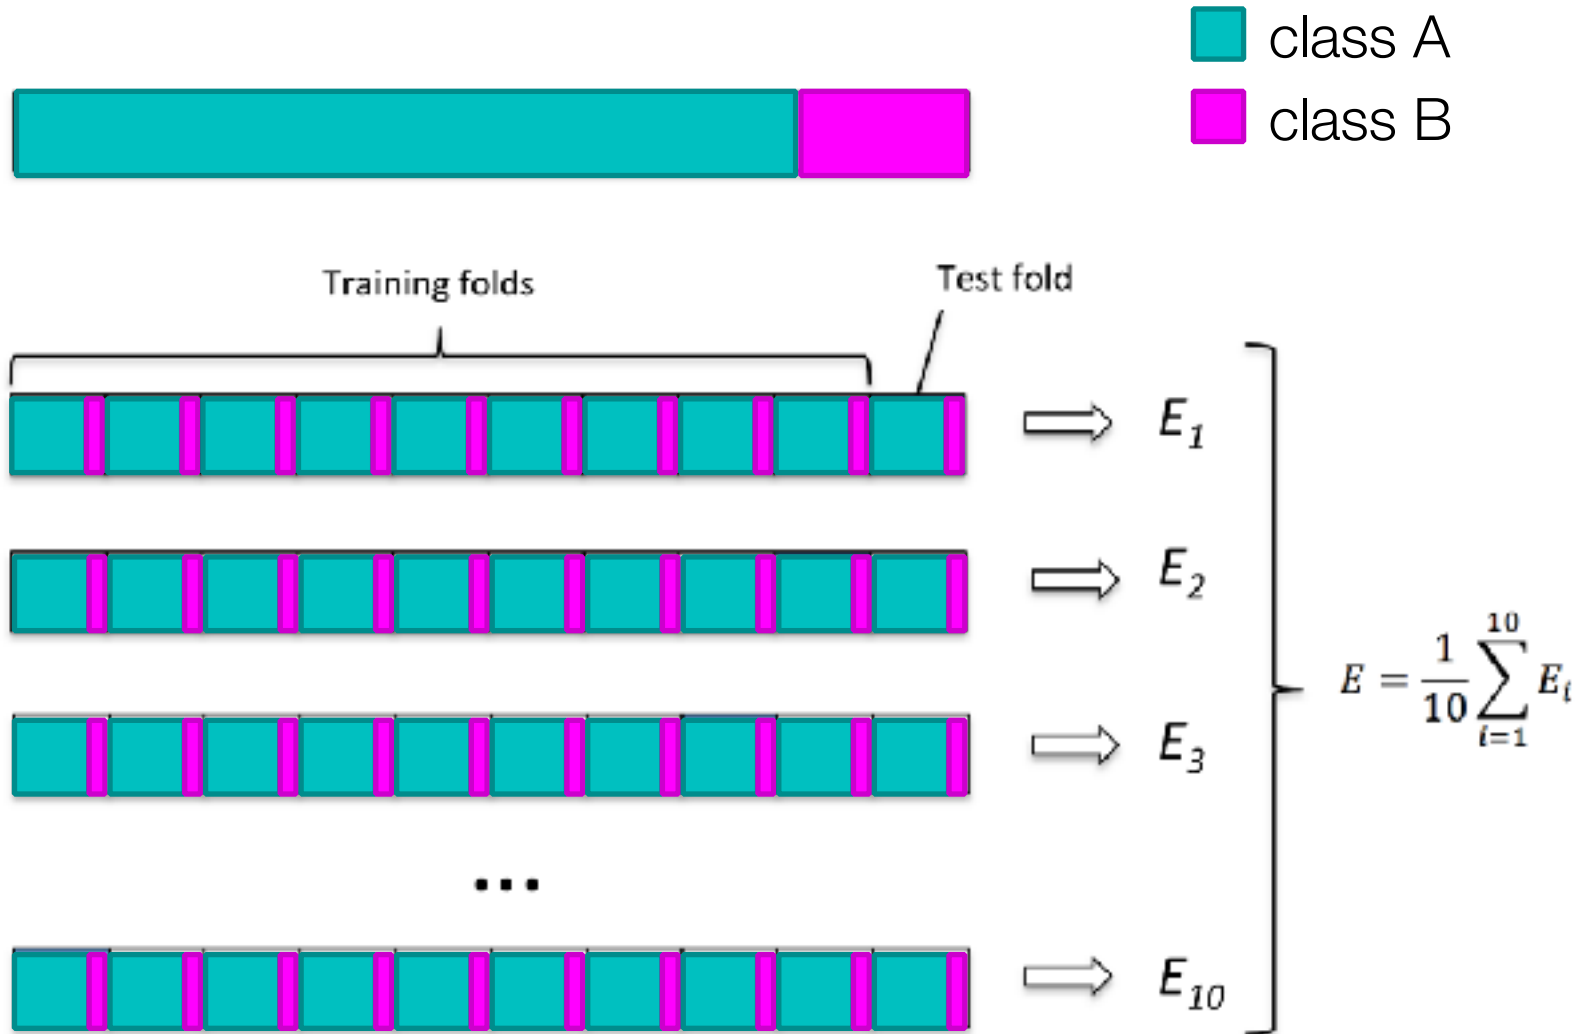
## $k$ -Fold Cross Validation





# Validation Loop Strategies

## Stratified $k$ -Fold Cross Validation

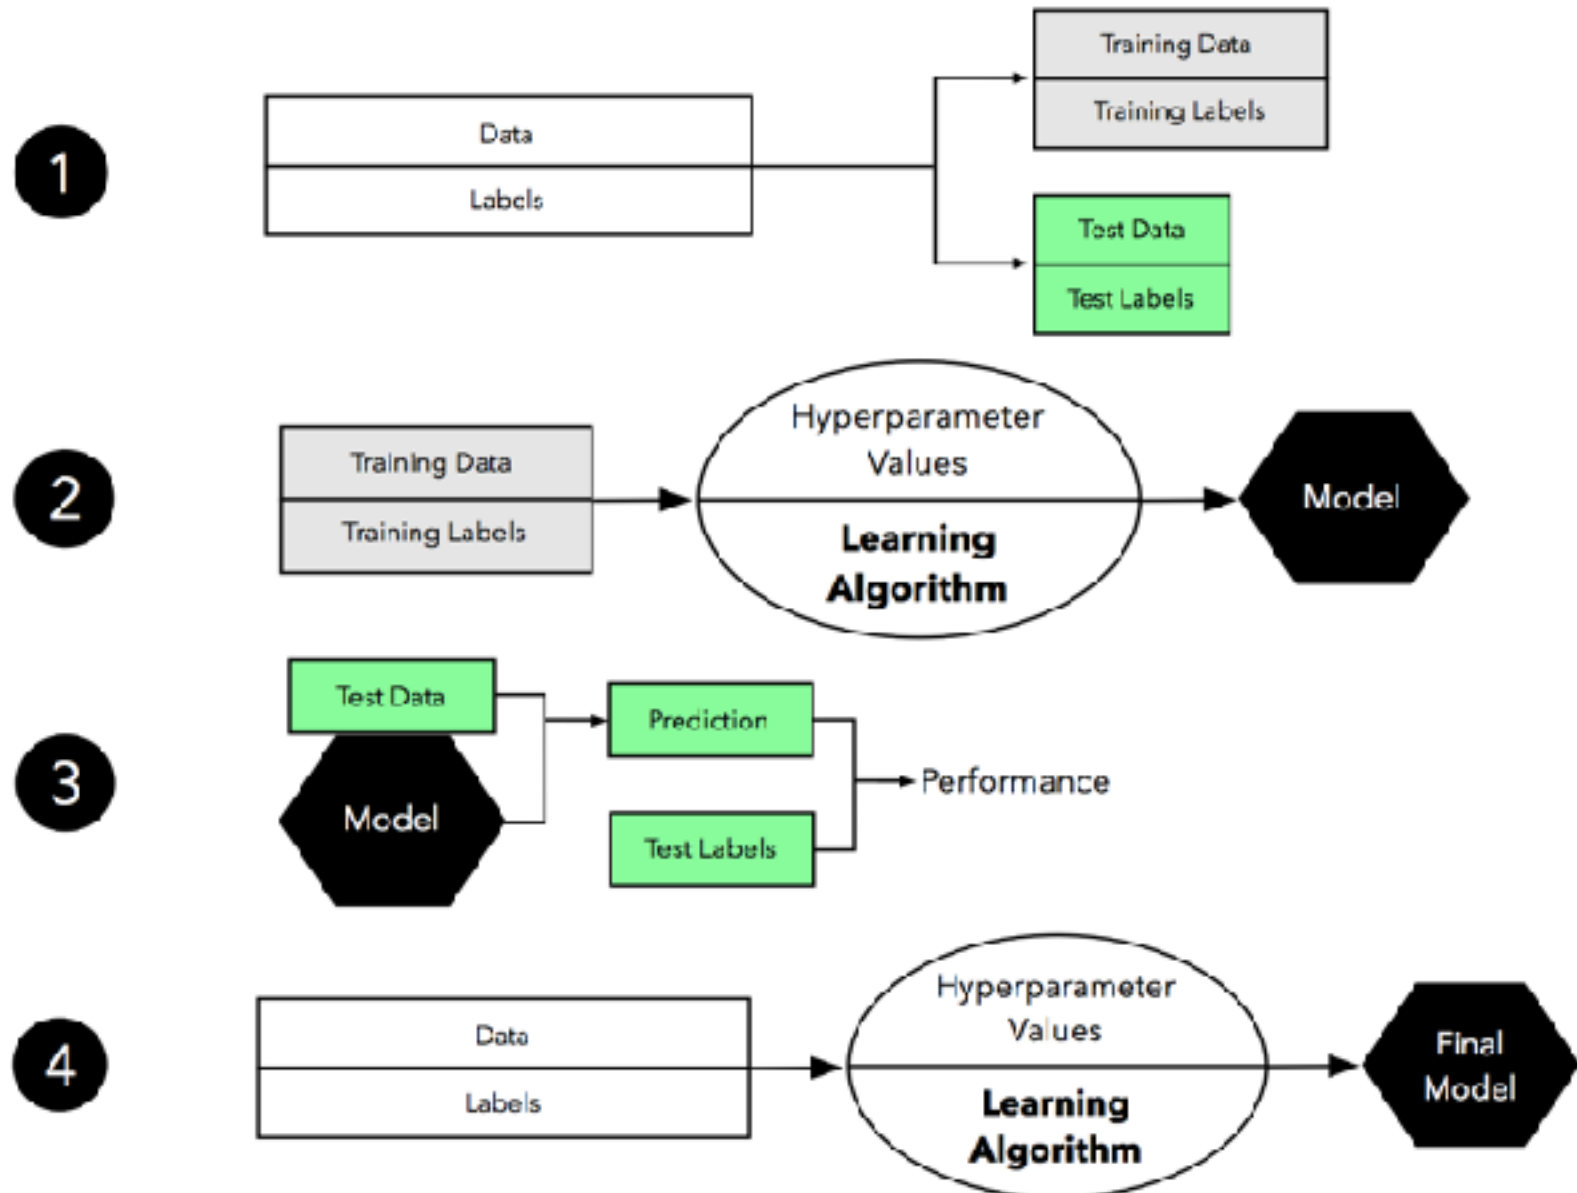


# Cross Validation: Reality Check

---


- What is the point of cross validation?
  - **Primary:** To realistically estimate how your classifier will perform on data it has not seen
    - situation must be ***plausible*** scenario
  - **Secondary:** hyperparameter tuning
    - what parameters should I set when training the model?
      - attain average performance from CV
      - use parameters to obtain final model

# Cross Validation: Reality Check




# Cross Validation: Reality Check

- Folds must be **plausible**, **representative** of the **actual use case** for the classifier
- Time series:  
**Cannot** apply stratified cross validation




# Cross Validation: Reality Check

- Folds must be **plausible**, **representative** of the **actual use case** for the classifier
- Time series:  
**Cannot** apply stratified cross validation
- Cannot** apply folding, actually



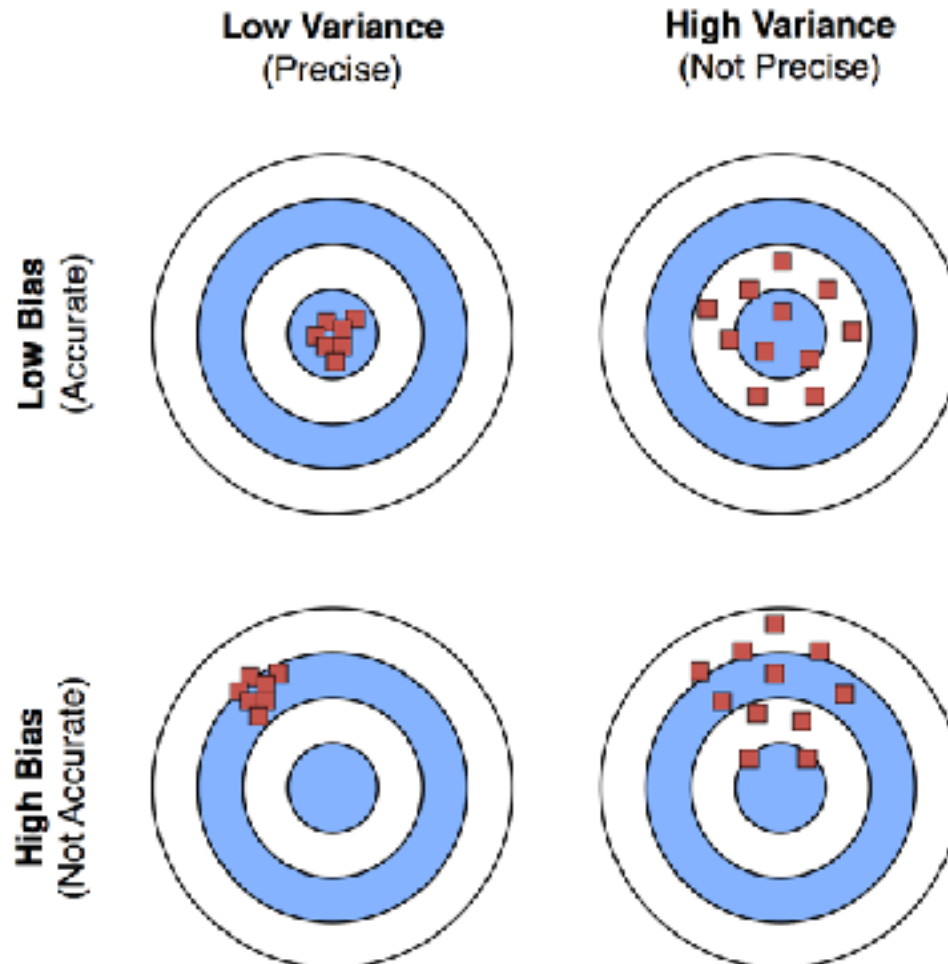
# Cross Validation: Reality Check

- Folds must be **plausible**, **representative** of the **actual use case** for the classifier
- Time series:  
**Cannot** apply stratified cross validation
- **Cannot** apply folding, actually
- Even better:  
**Mirror** real life use case




# Beyond cross validation

- So ... we separated out the data
- How do we know a model is actually good?



# Bias Variance Tradeoff

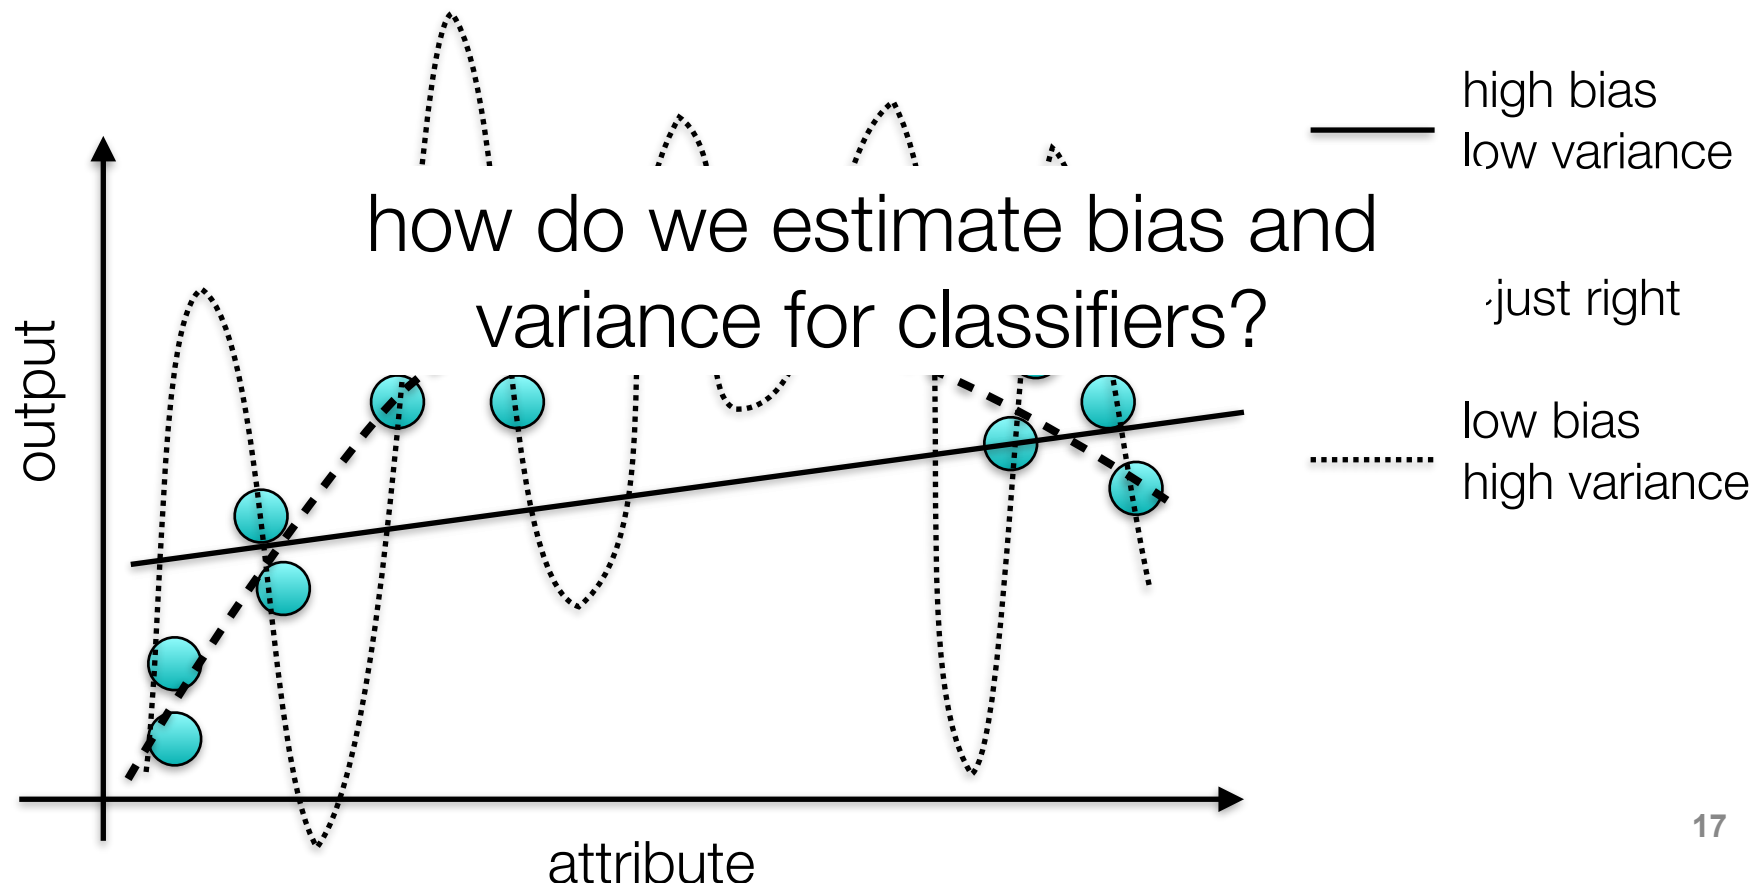
- **Complex** models can really fit the training data, giving **lower bias**
- **Simpler** models have trouble fitting data, resulting in **higher bias**
- But complex models can have **high variance** in their decision!






# Bias Variance Tradeoff

- **Complex** models can really fit the training data, giving **lower bias**
- **Simpler** models have trouble fitting data, resulting in **higher bias**
- But complex models can have **high variance** in their decision!



# The Learning Curve: Number of Samples



Learning curve shows how accuracy changes with varying sample size

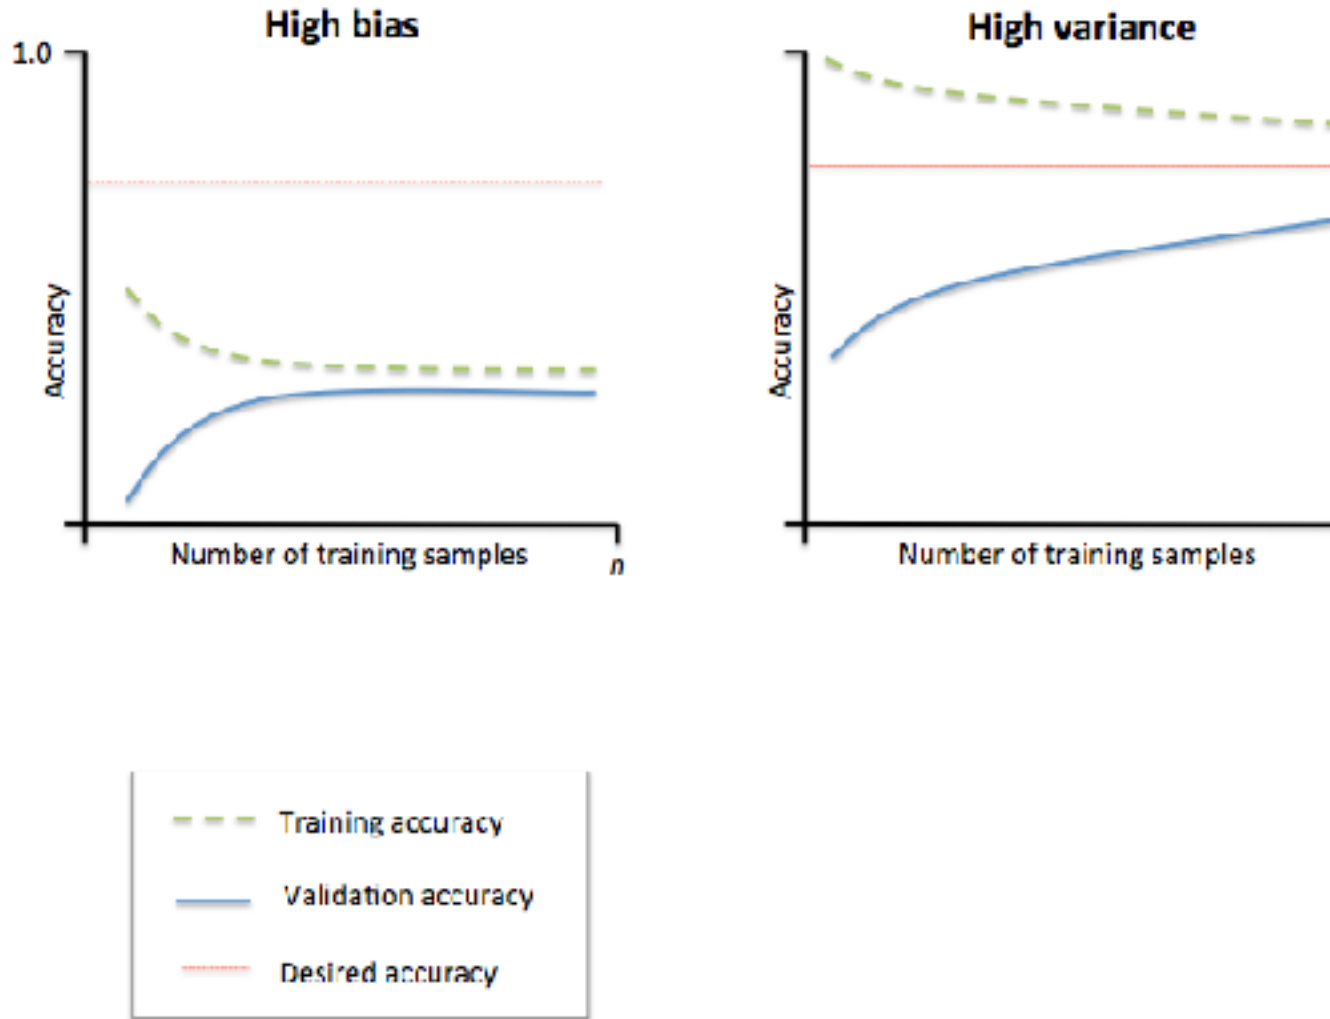
Effect of small sample size:

- Bias in the estimate
- Variance of estimate

You cannot estimate this curve without **collecting the data**. Some **bounds exist**, but they are **too loose** to be **useful!!!!**


randomly get this number of samples ...

# The Learning Curve: Number of Samples



# The Learning Curve: Number of Samples

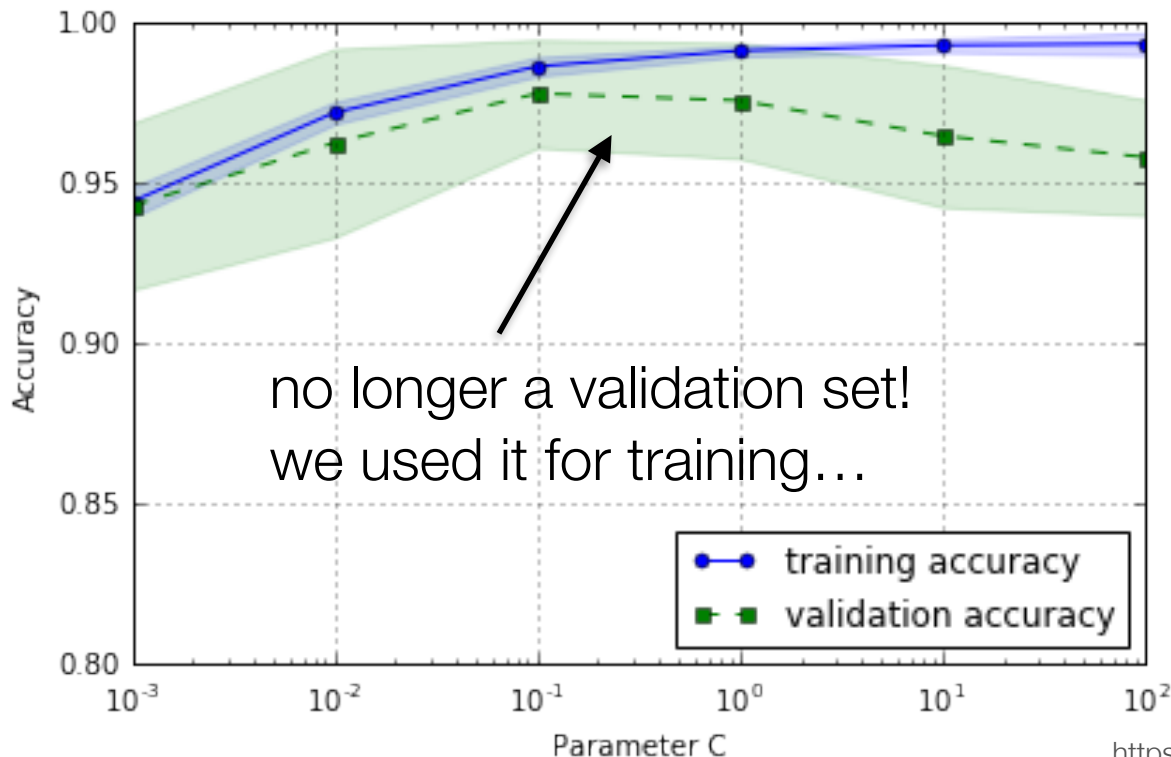
```
pipe_lr = Pipeline([('scl', StandardScaler()),  
                    ('clf', LogisticRegression(penalty='l2', random_state=0))])
```



# The Validation Curve: Hyper-parameters

```
param_range = [0.001, 0.01, 0.1, 1.0, 10.0, 100.0]
train_scores, test_scores = validation_curve(
    estimator=pipe_lr,
    X=X_train,
    y=y_train,
    param_name='clf__C',
    param_range=param_range,
    cv=10)
```

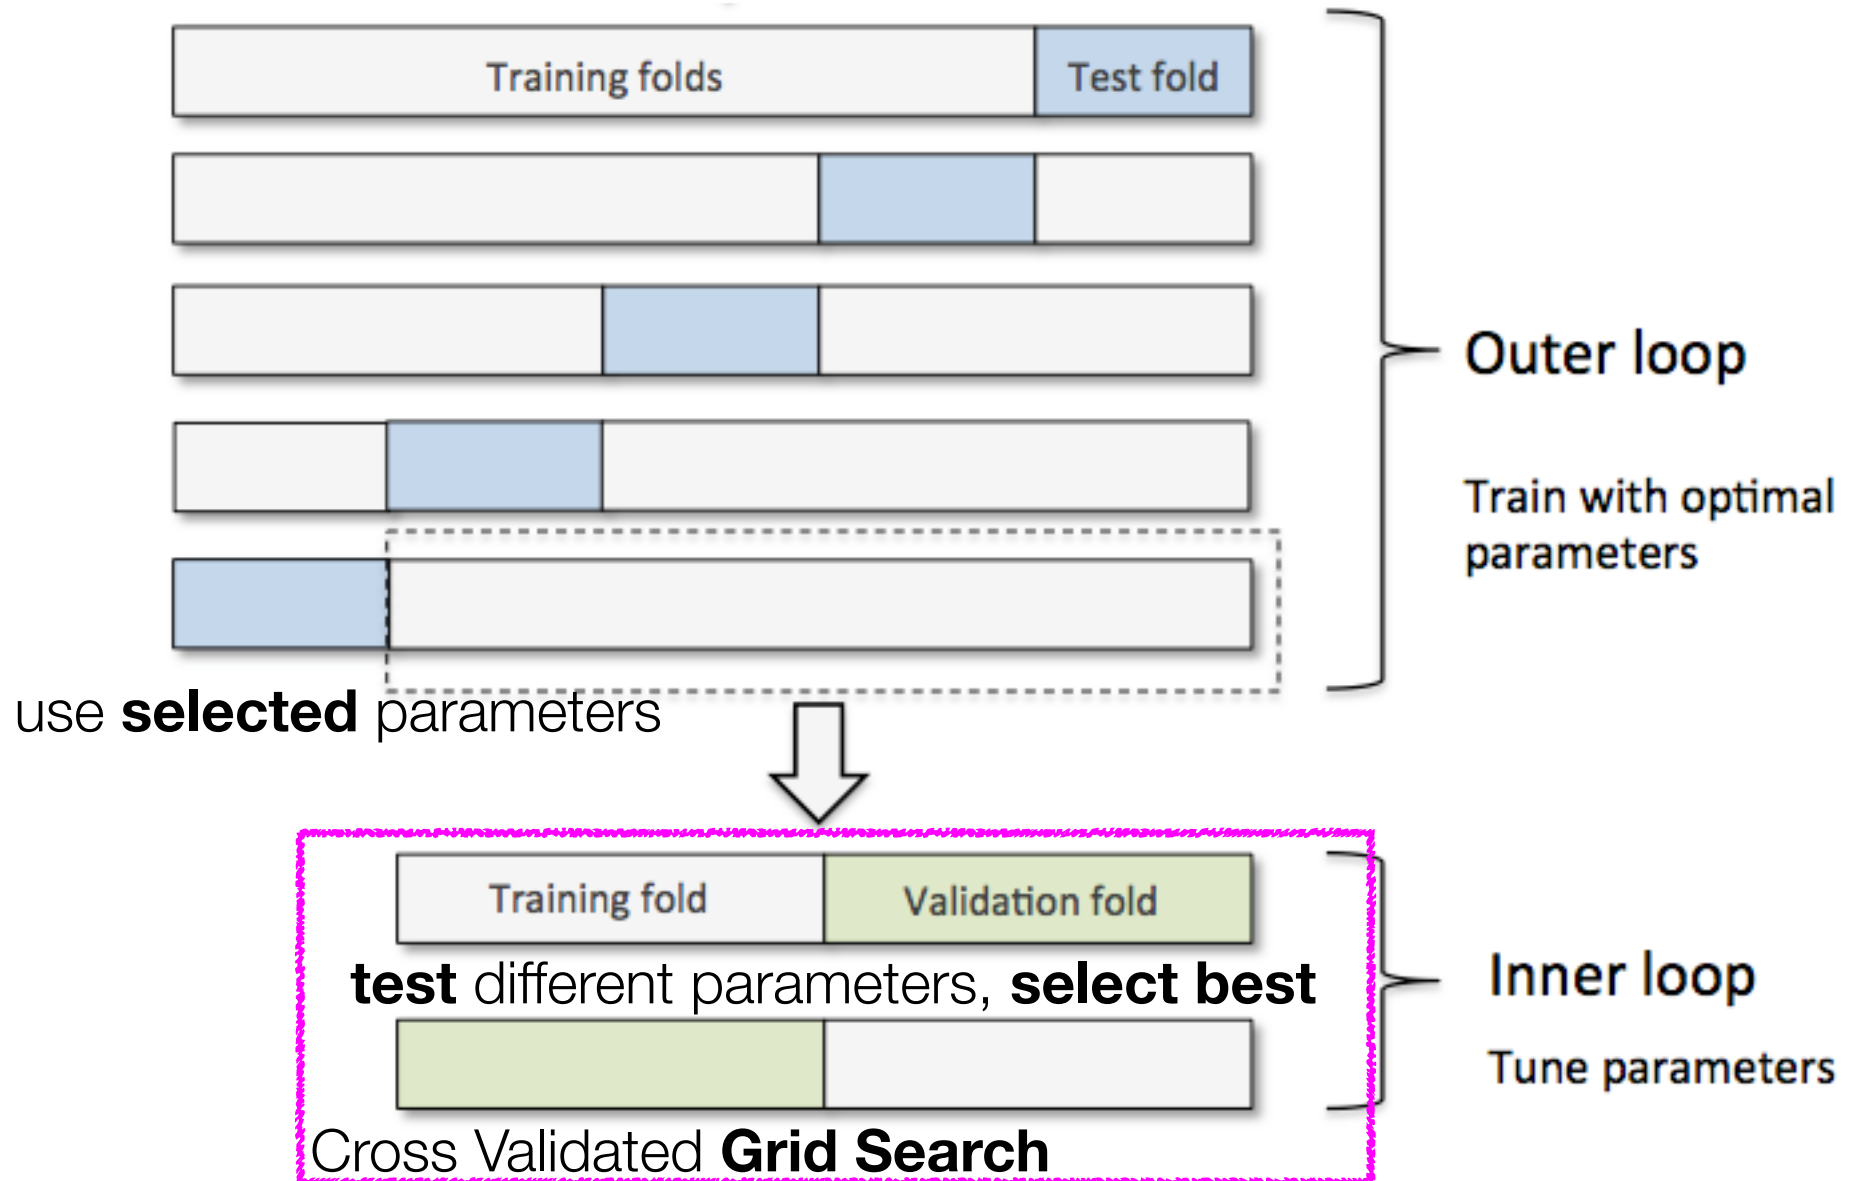
Similar to learning curve, but sweeping a hyper-parameter of the learning algorithm



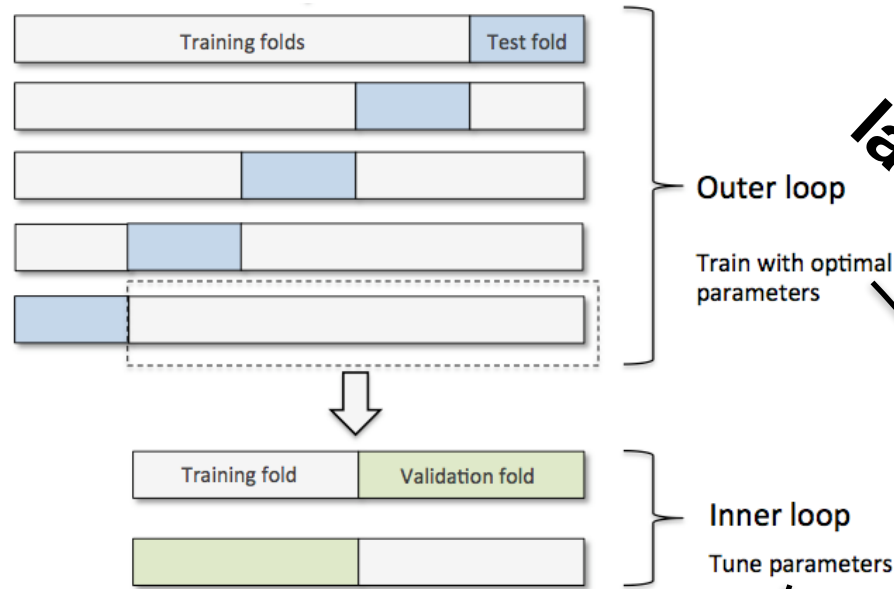
Look for sweet spot in the parameter for the validation accuracy

But, this can introduce data snooping...

# Nested Cross Validation: Hyper-parameters



# Nested Cross Validation: Hyper-parameters



We will **revisit** this  
**later** in semester

```
gs = GridSearchCV(estimator=pipe_svc,  
                  param_grid=param_grid,  
                  scoring='accuracy',  
                  cv=2)
```

*# Note: Optionally, you could use cv=2  
# in the GridSearchCV above to produce  
# the 5 x 2 nested CV that is shown in the figure.*

```
scores = cross_val_score(gs, X_train, y_train, scoring='accuracy', cv=5)  
print('CV accuracy: %.3f +/- %.3f' % (np.mean(scores), np.std(scores)))
```

---

# Model Evaluation Metrics/Measures

---

---



# Metrics for Performance Evaluation

- Focus on the **predictive capability** of a model
- **Not** how fast it takes to classify or build models, scalability, etc.
- Confusion Matrix:

|              | PREDICTED CLASS |          |           |
|--------------|-----------------|----------|-----------|
|              |                 | Class=No | Class=Yes |
| ACTUAL CLASS | Class=No        | d        | c         |
|              | Class=Yes       | b        | a         |

**a: TP (true positive)**

**b: FN (false negative)**

**c: FP (false positive)**

**d: TN (true negative)**

# Metrics for Performance Evaluation...

|              | PREDICTED CLASS |           |           |
|--------------|-----------------|-----------|-----------|
|              |                 | Class=No  | Class=Yes |
|              | Class=No        | d<br>(TN) | c<br>(FP) |
| ACTUAL CLASS | Class=Yes       | b<br>(FN) | a<br>(TP) |

- Most widely-used metric:

$$\text{Accuracy} = \frac{a + d}{a + b + c + d} = \frac{TP + TN}{TP + TN + FP + FN}$$

# Limitations of Accuracy

---

- Ignores the **cost** of misclassifications
- Consider an **imbalanced** 2-class problem
  - Number of Class 0 examples = 9990
  - Number of Class 1 examples = 10
- If model **predicts everything to be class 0**, accuracy is  $9990/10000 = 99.9\%$ 
  - Accuracy is **misleading** because model does not detect any class 1 example

# Other evaluation metrics: Cost Matrix

|              | PREDICTED CLASS |                           |                            |
|--------------|-----------------|---------------------------|----------------------------|
| ACTUAL CLASS | $C(i   j)$      | Class=No                  | Class=Yes                  |
|              | Class=No        | $C(\text{No} \text{No})$  | $C(\text{Yes} \text{No})$  |
|              | Class=Yes       | $C(\text{No} \text{Yes})$ | $C(\text{Yes} \text{Yes})$ |

Define a cost function based on your expertise with problem:

$C(i | j)$ : Cost of misclassifying class  $j$  example as class  $i$

# Cost Matrix Examples

Lower cost  
means “better”

| Cost Matrix  | PREDICTED CLASS |     |    |
|--------------|-----------------|-----|----|
| ACTUAL CLASS | C(i j)          | -   | +  |
|              | -               | 0   | 1  |
|              | +               | 100 | -1 |

*i.e.*, medical  
diagnosis costs?

| Cost Matrix  | PREDICTED CLASS |     |     |     |    |
|--------------|-----------------|-----|-----|-----|----|
| ACTUAL CLASS | C(i j)          | USA | CAN | AUS | NZ |
|              | USA             | 0   | 1   | 10  | 10 |
|              | CAN             | 1   | 0   | 10  | 10 |
|              | AUS             | 10  | 10  | 0   | 20 |
|              | NZ              | 10  | 10  | 20  | 0  |

*i.e.*, predicting  
travel  
locations?

# Cost-Sensitive Measures

|  | PREDICTED CLASS |
|--|-----------------|
|--|-----------------|

```
from sklearn.metrics import precision_score, recall_score, f1_score

print('Precision: %.3f' % precision_score(y_true=y_test, y_pred=y_pred))
print('Recall: %.3f' % recall_score(y_true=y_test, y_pred=y_pred))
print('F1: %.3f' % f1_score(y_true=y_test, y_pred=y_pred))
```

| CLASS |           | (FP)      | (FN)      |
|-------|-----------|-----------|-----------|
|       | Class=Yes | b<br>(FN) | a<br>(TP) |

$$\text{Precision (p)} = \frac{a}{a + c}$$

Higher Precision ==  
Lower false positives

$$\text{Recall (r)} = \frac{a}{a + b}$$

Higher Recall ==  
Lower false negatives

$$\text{F - measure (F)} = \frac{2rp}{r + p} = \frac{2a}{2a + b + c}$$

Higher F1 ==  
Lower FN & FP

# Cost-Sensitive Measures: Multi-Class

```
pre_scorer = make_scorer(score_func=precision_score,  
                          pos_label=1,  
                          greater_is_better=True,  
                          average='micro')
```

$$\text{Precision}_{\text{micro}} = \frac{TP_1 + \dots + TP_k}{TP_1 + \dots + TP_k + FP_1 + \dots + FP_k}$$

$$\text{Recall}_{\text{micro}} = \frac{TP_1 + \dots + TP_k}{TP_1 + \dots + TP_k + FN_1 + \dots + FN_k}$$

weight all instances equally

$$F1_{\text{micro}} = \frac{2(TP_1 + \dots + TP_k)}{2(TP_1 + \dots + TP_k) + FP_1 + \dots + FP_k + FN_1 + \dots + FN_k}$$

$$X_{\text{macro}} = \frac{X_1 + \dots + X_k}{k}$$

weight all classes equally

---

# The Receiver Operating Characteristic

---

---



# ROC (Receiver Operating Characteristic)

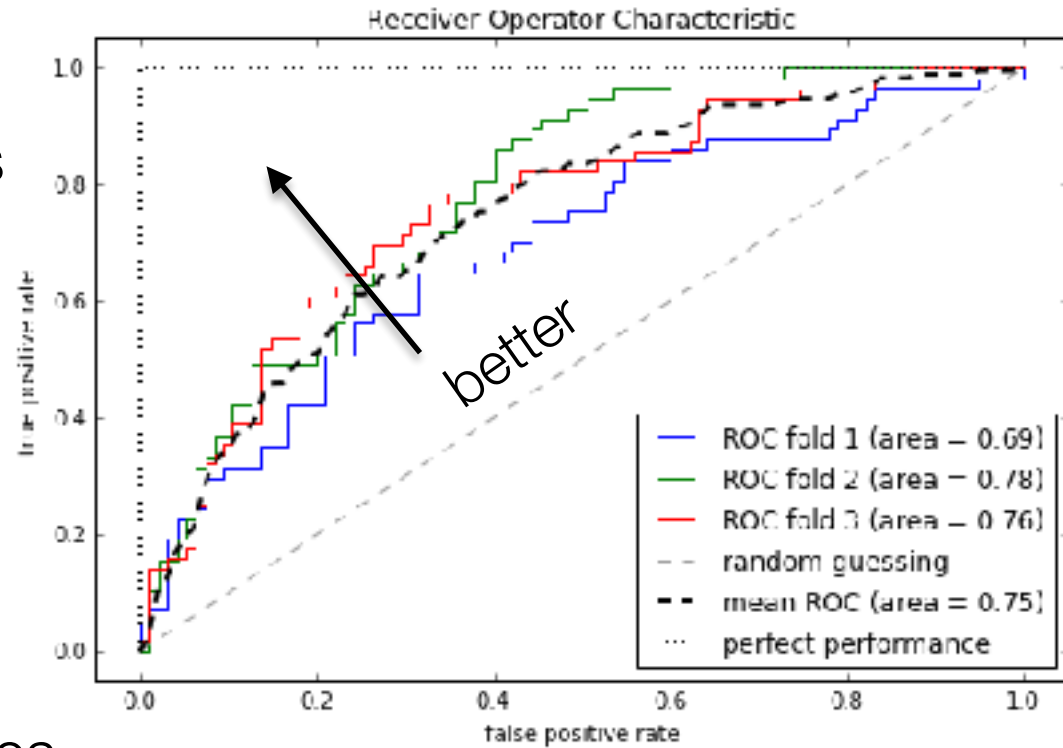
---

- Developed in 1950s for signal detection theory to analyze noisy signals
  - Characterize the trade-off between positive hits and false alarms
- ROC curve plots TP (on the y-axis) against FP (on the x-axis)
- Performance of each classifier represented as a point on the ROC curve
  - changing the threshold of algorithm, sample distribution or cost matrix changes the location of the point


# ROC Curve

(TP,FP):

- (0,0): declare everything to be negative class
- (1,1): declare everything to be positive class
- (1,0): ideal
- Diagonal line:
  - Random guessing for equal number of classes
  - Below diagonal line:
    - ♦ prediction is opposite of the true class



# Using ROC for Model Comparison



- No model consistently outperforms the other
- $M_1$  is better for small FPR
- $M_2$  is better for large FPR
- Area Under the ROC curve
- Ideal: Area = 1.0

# How to Construct an ROC curve

classifier score

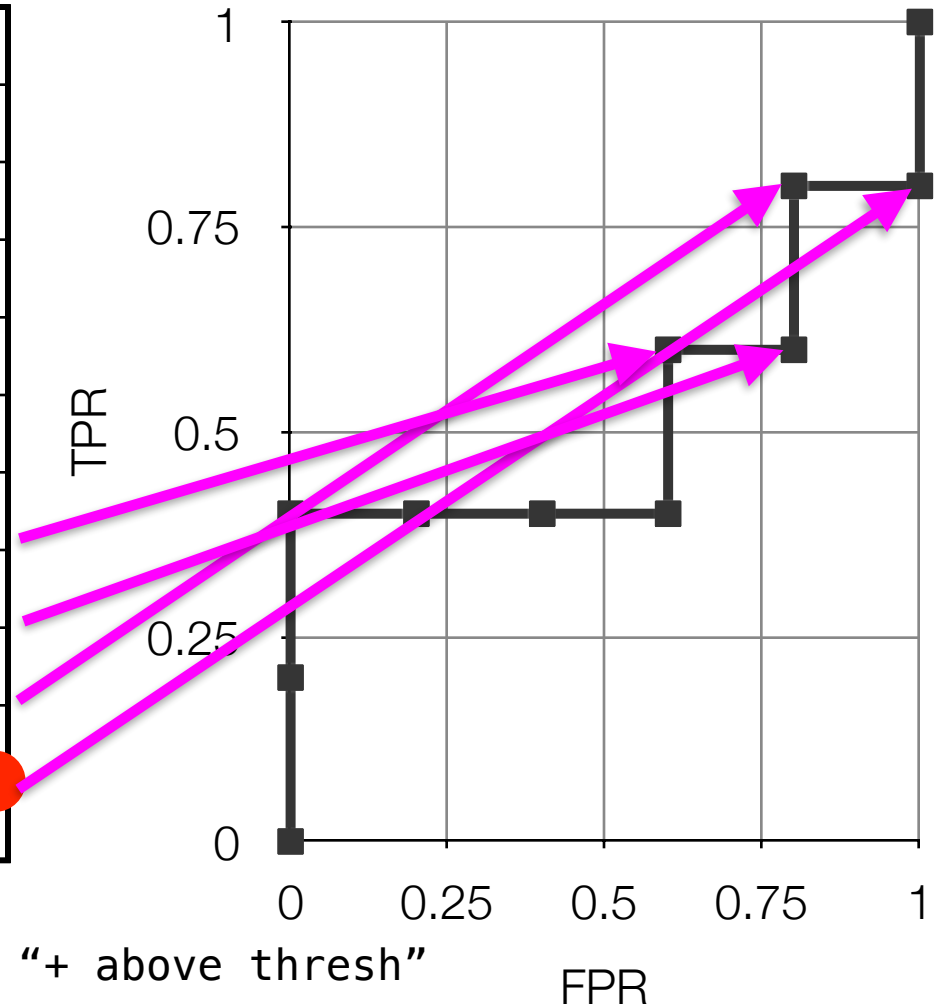
| Instance # | $P(+ A)$ | True Class |
|------------|----------|------------|
| 6          | 0.95     | +          |
| 2          | 0.93     | +          |
| 5          | 0.87     | -          |
| 4          | 0.85     | -          |
| 9          | 0.85     | -          |
| 1          | 0.85     | +          |
| 10         | 0.76     | -          |
| 8          | 0.53     | +          |
| 3          | 0.43     | -          |
| 7          | 0.25     | +          |

- Use classifier that produces probability score for each test instance  $P(+|A)$
- Sort the instances according to  $P(+|A)$  in decreasing order
- Apply threshold,  $T$ , at each unique value of  $P(+|A)$
- $P(+|A) < T$ , is negative class, else it is a positive class
- Count the number of TP, FP, TN, FN at each threshold
- TP rate,  $TPR = TP/Positives$
- FP rate,  $FPR = FP/Negatives$ <sup>36</sup>

# How to Construct an ROC curve

classifier score

| Instance # | P(+ A) | True Class |
|------------|--------|------------|
| 6          | 0.95   | +          |
| 2          | 0.93   | +          |
| 5          | 0.87   | -          |
| 4          | 0.85   | -          |
| 9          | 0.85   | -          |
| 1          | 0.85   | +          |
| 10         | 0.76   | -          |
| 8          | 0.53   | +          |
| 3          | 0.43   | -          |
| 7          | 0.25   | +          |



- TP rate,  $TPR = TP / \text{Positives}$ , “+ above thresh”
- FP rate,  $FPR = FP / \text{Negatives}$  “- above thresh”

# How to Construct an ROC curve

---

```
for i, (train, test) in enumerate(cv):  
    probas = pipe_lr.fit(X_train2[train],  
                        y_train[train]).predict_proba(X_train2[test])  
  
    fpr, tpr, thresholds = roc_curve(y_train[test],  
                                    probas[:, 1],  
                                    pos_label=1)
```

---

# Significance Testing

---

---

# Tests of Significance

---

- Given two models:
  - Model  $M_1$ : accuracy = 85%, tested on 30 instances
  - Model  $M_2$ : accuracy = 75%, tested on 5000 instances
- Can we say  $M_1$  is better than  $M_2$ ?
  - How much confidence can we place on accuracy of  $M_1$  and  $M_2$ ?
  - Can the difference in performance measure be explained as a result of random fluctuations in the test set?



# Comparing Performance of 2 Models

- Given two models,  $M_1$  and  $M_2$ , which is better?
  - $M_1$  is tested on  $D_1$  (size= $n_1$ ), found error rate =  $e_1$
  - $M_2$  is tested on  $D_2$  (size= $n_2$ ), found error rate =  $e_2$
  - Assume  $D_1$  and  $D_2$  are independent
  - If  $n_1$  and  $n_2$  are sufficiently large, then

$$e_1 \sim N(\mu_1, \sigma_1)$$


$$e_2 \sim N(\mu_2, \sigma_2)$$

- Approximate: 
$$\hat{\sigma}_i^2 = \frac{e_i(1 - e_i)}{n_i}$$

variance estimate comes from **binomial distribution**,  
which is approximated well by **normal distribution**

# Comparing Performance of 2 Models

- To test if performance difference is statistically significant:  $d = e_1 - e_2$  ← **estimate of the mean difference**
  - $d \sim N(d_t, \sigma_t)$  where  $d_t$  is the true difference
  - Since  $D_1$  and  $D_2$  are independent, their variance adds up:

$$\begin{aligned}\sigma_t^2 &= \sigma_1^2 + \sigma_2^2 \cong \hat{\sigma}_1^2 + \hat{\sigma}_2^2 \\ &= \frac{e1(1-e1)}{n1} + \frac{e2(1-e2)}{n2}\end{aligned}$$


**estimate of the variance in subtracted error rates**

- At  $(1-\alpha)$  confidence level, bounds on  $d = d \pm Z_{\alpha/2} \hat{\sigma}_t$

**does this interval include zero?**

# An Illustrative Example

$$= d \pm Z_{\alpha/2} \hat{\sigma}_t$$

- Given:  $M_1: n_1 = 30, e_1 = 0.15$   
 $M_2: n_2 = 5000, e_2 = 0.25$
- $d = |e_2 - e_1| = 0.1$  (2-sided test)

$$\hat{\sigma}_d^2 = \frac{0.15(1-0.15)}{30} + \frac{0.25(1-0.25)}{5000} = 0.0043$$

- At 95% confidence level,  $Z_{\alpha/2} = 1.96$

$$d_t = 0.100 \pm 1.96 \times \sqrt{0.0043} = 0.100 \pm 0.128$$

- $\Rightarrow$  Interval contains 0  $\Rightarrow$  difference may not be statistically significant

# Another illustrative example

---



**“THERE ARE LIES, DAMNED LIES AND  
STATISTICS.”**

**MARK TWAIN**

© Lifhack Quotes

$M_1: n_1 = 30, e_1 = 0.15$

$M_2: n_2 = 5000, e_2 = 0.25$

Use common sense and choose the classifier with 5000 test examples

# Folded statistical comparisons

- Each learning algorithm may produce  $k$  models:
  - $L_1$  may produce  $M_{11}, M_{12}, \dots, M_{1k}$
  - $L_2$  may produce  $M_{21}, M_{22}, \dots, M_{2k}$
- If models are generated on the same test sets  $D_1, D_2, \dots, D_k$  (e.g., via cross-validation)
  - For each set: compute  $d_j = e_{1j} - e_{2j}$ , the  $j^{\text{th}}$  difference
  - $d_j$  has mean  $\bar{d}$  and variance  $\sigma_t^2$

$$\sigma_t^2 = \frac{1}{k-1} \sum_{j=1}^k (\bar{d} - d_j)^2$$

$$d_t = \bar{d} \pm \frac{1}{\sqrt{k}} t_{1-\alpha, k-1} \sigma_t$$

mean of folds

confidence  
multiplier

now we can bound to  
get a better idea about  
how the criterion varies

$$t(95\%, k=10) = 2.26$$

standard deviation of folds

# An illustrative Example

```
acc1 = cross_val_score(clf1, X, y=y, cv=cv)
acc2 = cross_val_score(clf2, X, y=y, cv=cv)
```

```
#=====
```

```
t = 2.26 / np.sqrt(10)
```

```
e = (1-acc1)-(1-acc2)
```

```
# std1 = np.std(acc1)
```

```
# std2 = np.std(acc2)
```

```
stdtot = np.std(e)
```

```
dbar = np.mean(e)
```

```
print 'Range of:', dbar-t*stdtot, dbar+t*stdtot
```

$$\sigma_t^2 = \frac{1}{k-1} \sum_j^k (\bar{d} - d_j)^2$$
$$d_t = \bar{d} \pm \frac{1}{\sqrt{k}} t_{1-\alpha, k-1} \sigma_t$$

---

## **For in class assignment:**

**Using Cross Validation, Folding,  
ROC, and Statistical Significance  
to Evaluate Algorithm  
Performance**